



OpenVINO™
DEVCON
Workshop Series 2024

Model Development and Deployment with the OpenVINO™ Ecosystem

Our speakers



Paula Ramos
AI Software Evangelist



Samet Akcay
AI Research Engineer
& Scientist

Today

01
**Computer Vision
Solutions**

02
OpenVINO Toolkit

03
Anomalib

04
CVPR Challenge

Computer Vision Solutions

AI is Everywhere – from Edge to Cloud



Edge

- Real time data processing
- Wider reach
- Data sovereignty
- Cost efficiency



Frictionless Retail



Traffic Monitoring



Defect Detection



Choose Your Compute

Light AI

Efficiency for sub-100W designs
CPU AI, built-in GPU, built-in NPU

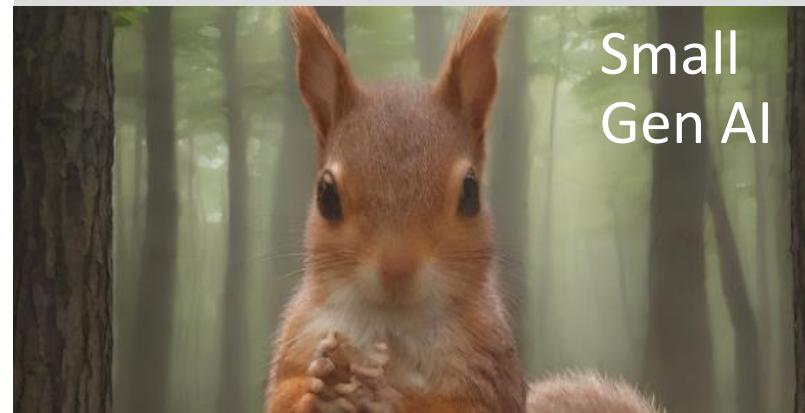


Medium AI

Scale up perf/W for diverse system designs
Discrete GPU & built-in CPU AI



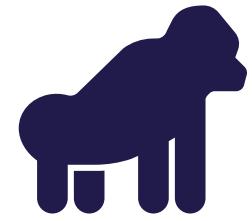
Small
Gen AI



Heavy AI

Optimize for peak perf and density
Optimize with high-end Discrete GPU





Large
model size



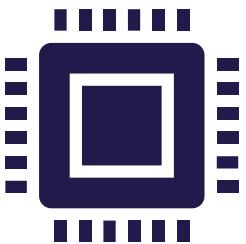
Large memory
footprint



Slow inference
speed

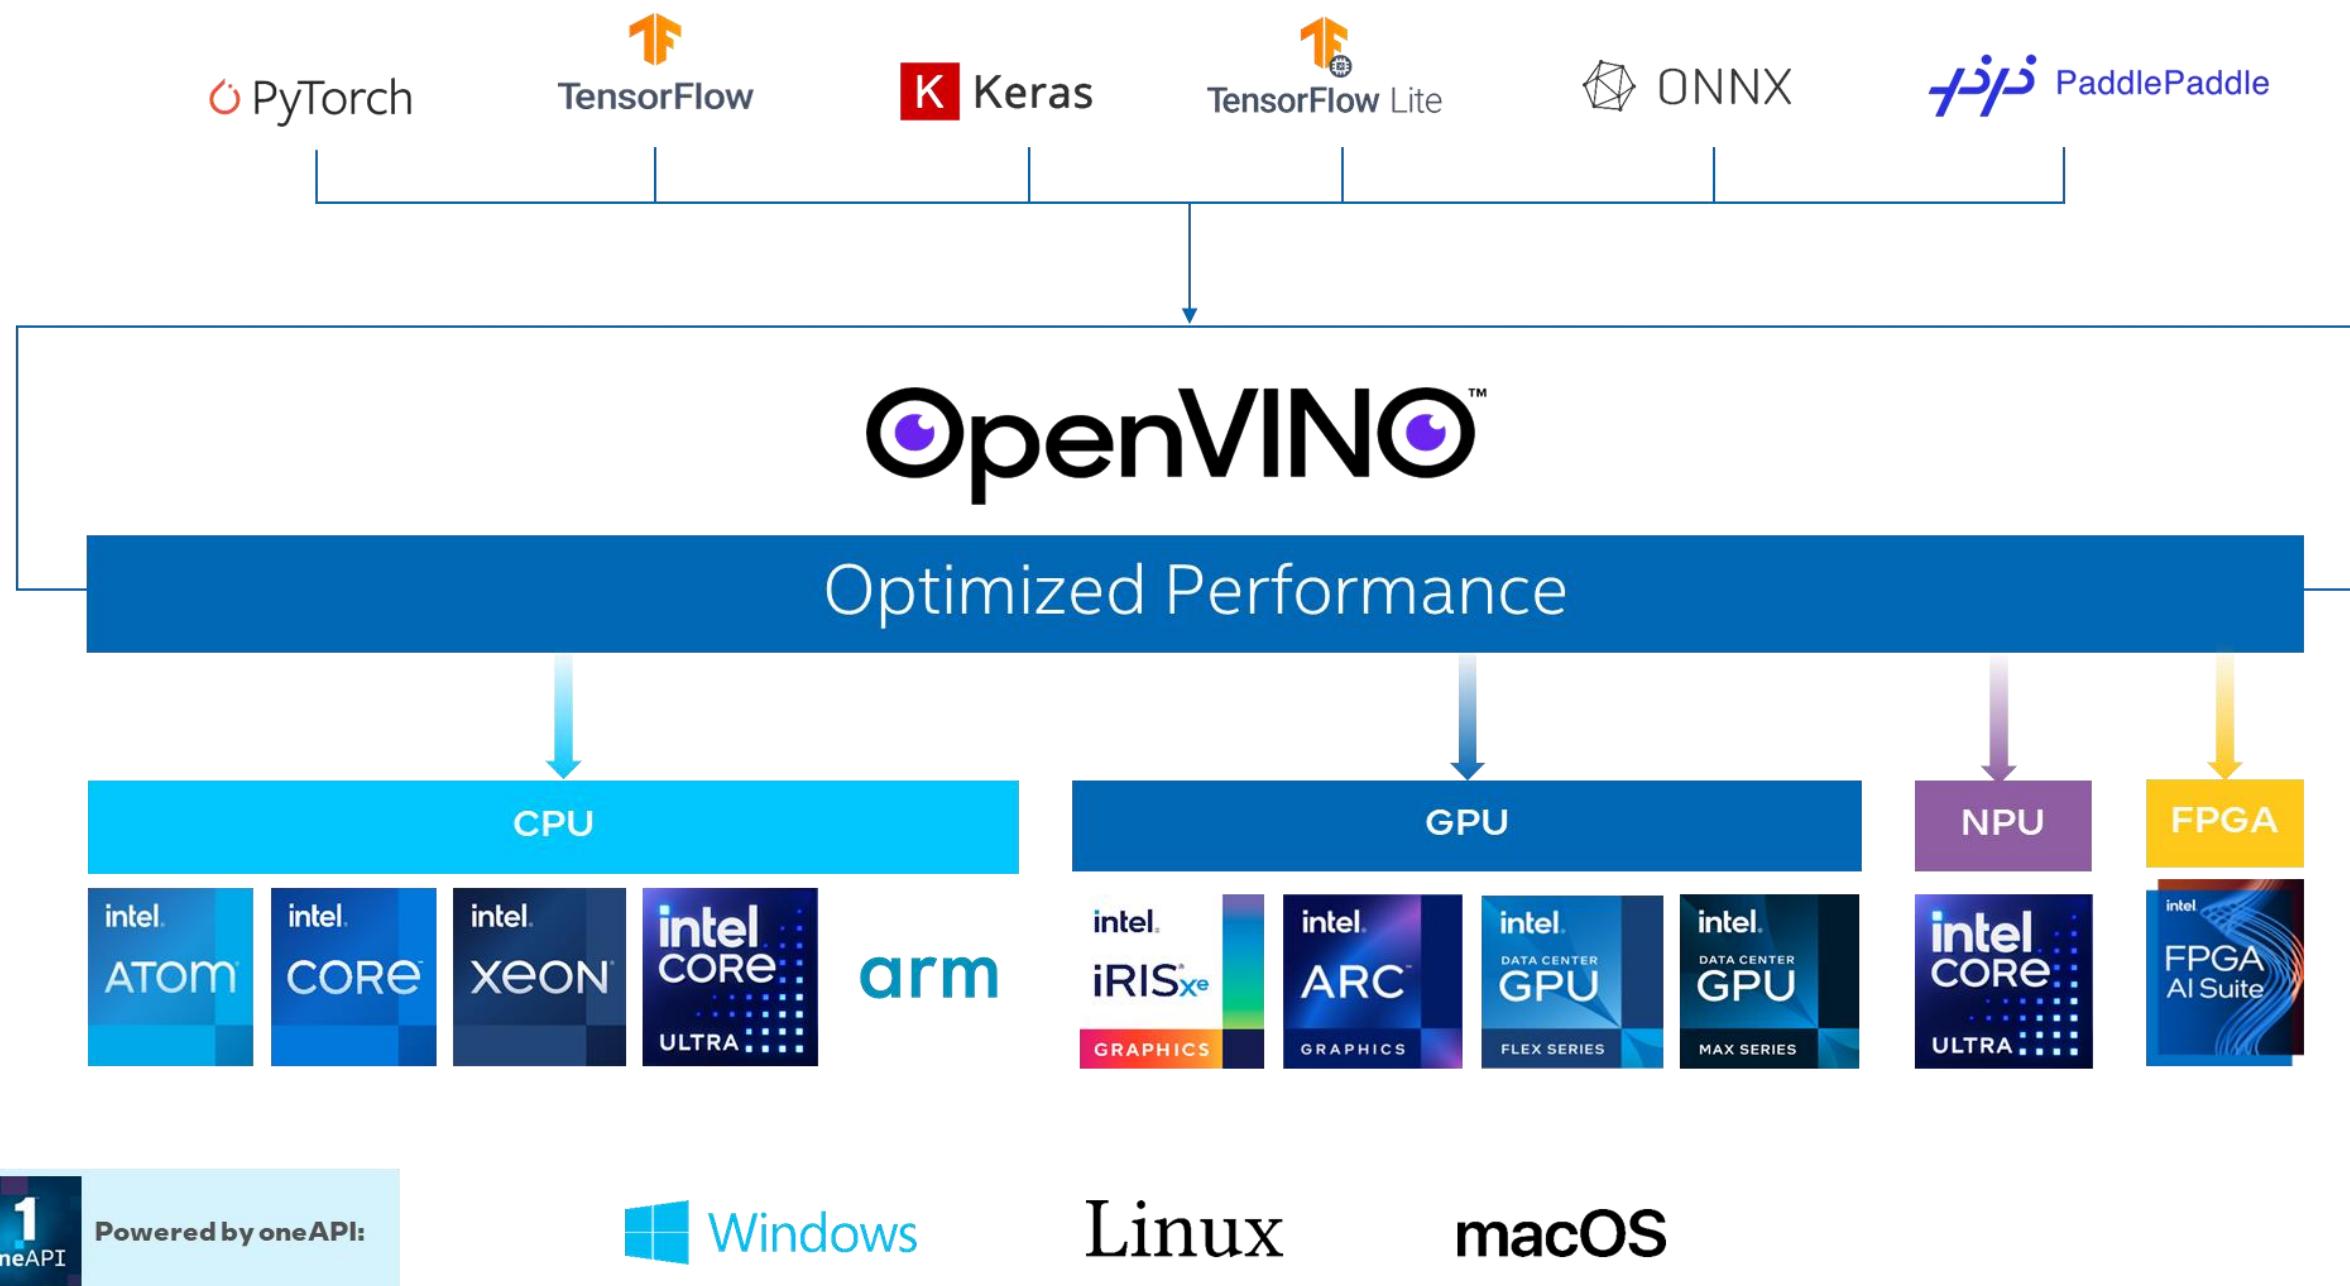


Difficulty training
+ optimizing

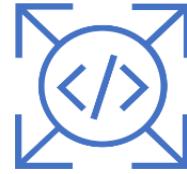


No flexibility to
run workloads
on different HW

An Open Standard for Building AI at the Edge



Benefits of Building Applications with OpenVINO™



Build and
deploy AI
applications in
simple steps



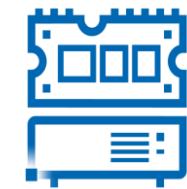
Faster
inference speed



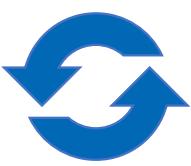
Maximize AI
performance
across CPU,
GPU, NPU



Smaller
model & binary size

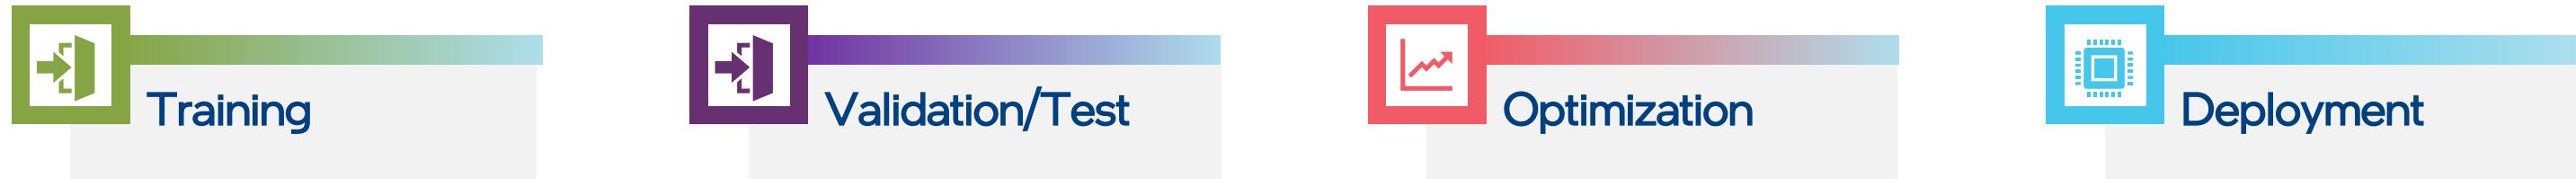


Reduce
memory footprint



Ability to scale
to many nodes
with serving

What is going on with the rest of the workflow?



OpenVINO™

Intel open-source library

Real-world datasets
are highly imbalanced

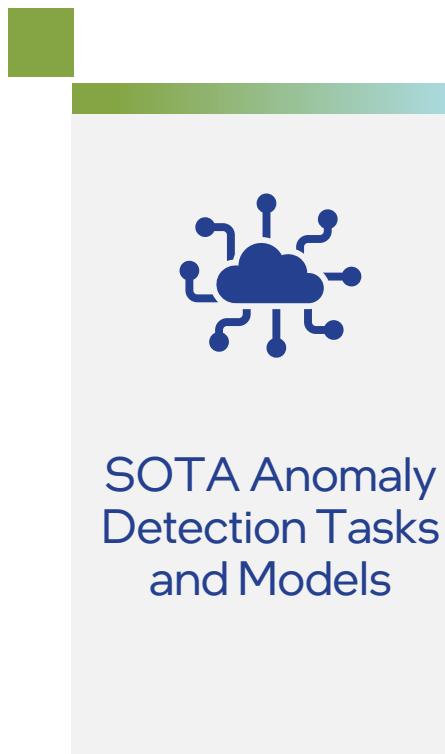
Abnormalities are often unknown
and might evolve/change

Anomaly Detection

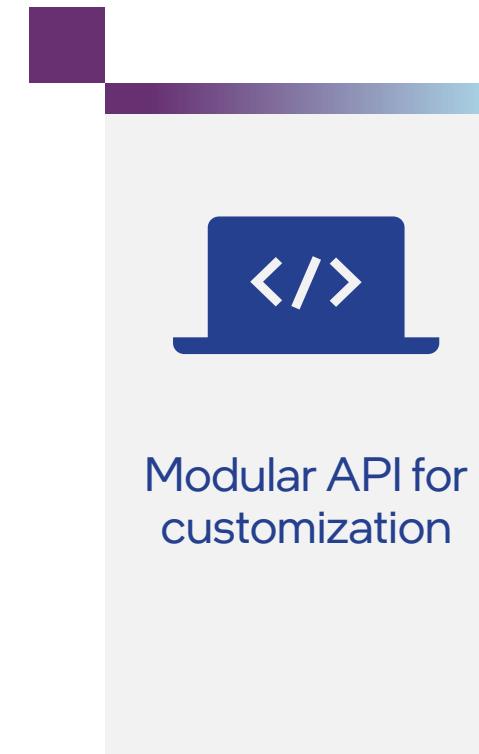
Anomalib

Anomalib is a library to
design, develop and deploy
DL anomaly detection algorithms

Key Features



SOTA Anomaly
Detection Tasks
and Models



Modular API for
customization

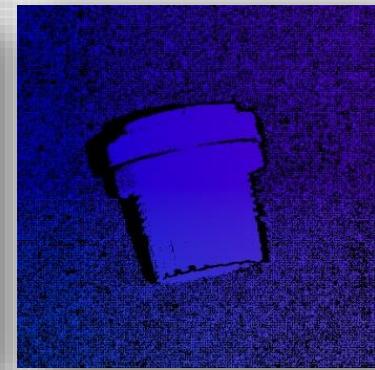
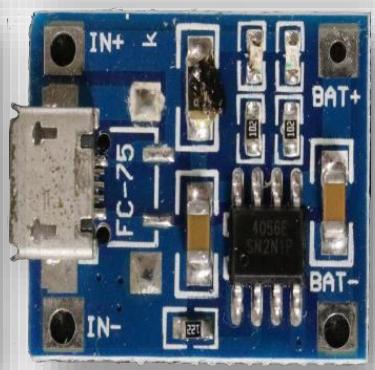
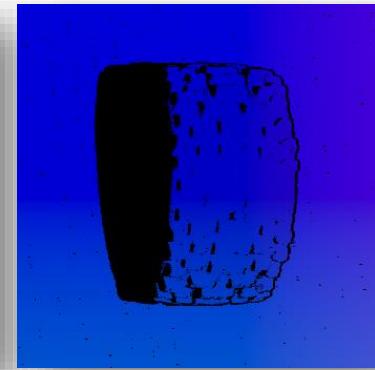


Fully
configurable CLI
for end-to-end
machine learning
life cycle



Various
inference
interfaces for
deployment

Supported Domains



Image

Depth

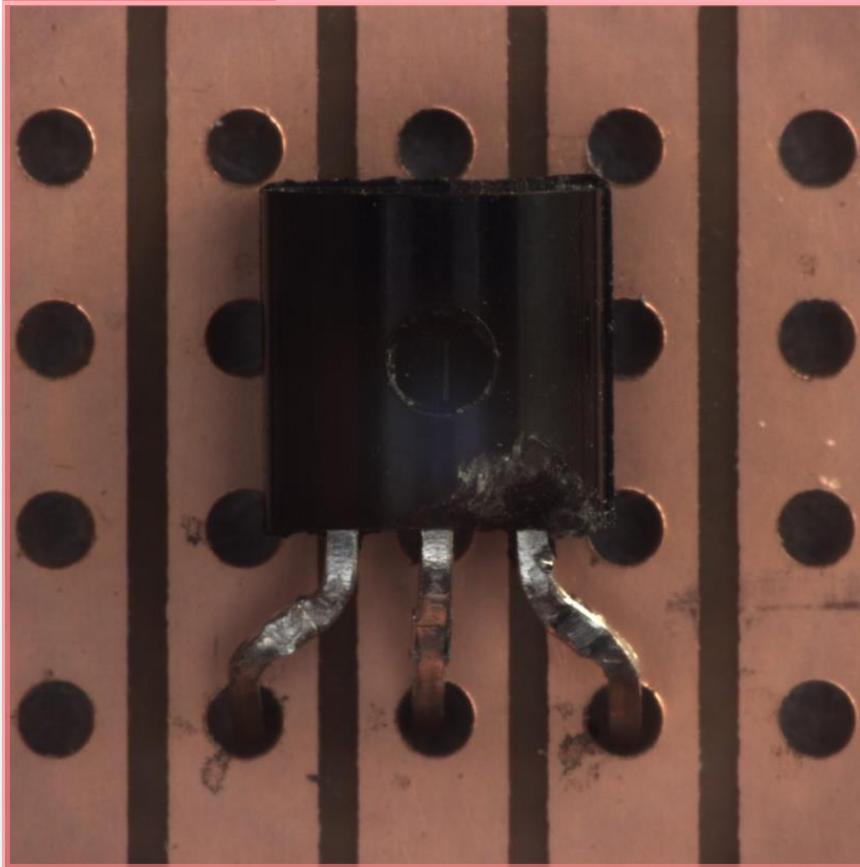
Video

P. Bergmann *et al.* The MVTec Anomaly Detection Dataset: A Comprehensive Real-World Dataset for Unsupervised Anomaly Detection, IJCV 2021
Zou, Y., Jeong, J., Pemula, L., Zhang, D., & Dabeer, O. SPot-the-Difference Self-supervised Pre-training for Anomaly Detection and Segmentation, ECCV 2022
P. Bergmann, X. Jin, D. Sattlegger, C. Steger: The MVTec 3D-AD Dataset for Unsupervised 3D Anomaly Detection and Localization, VISAPP 2022

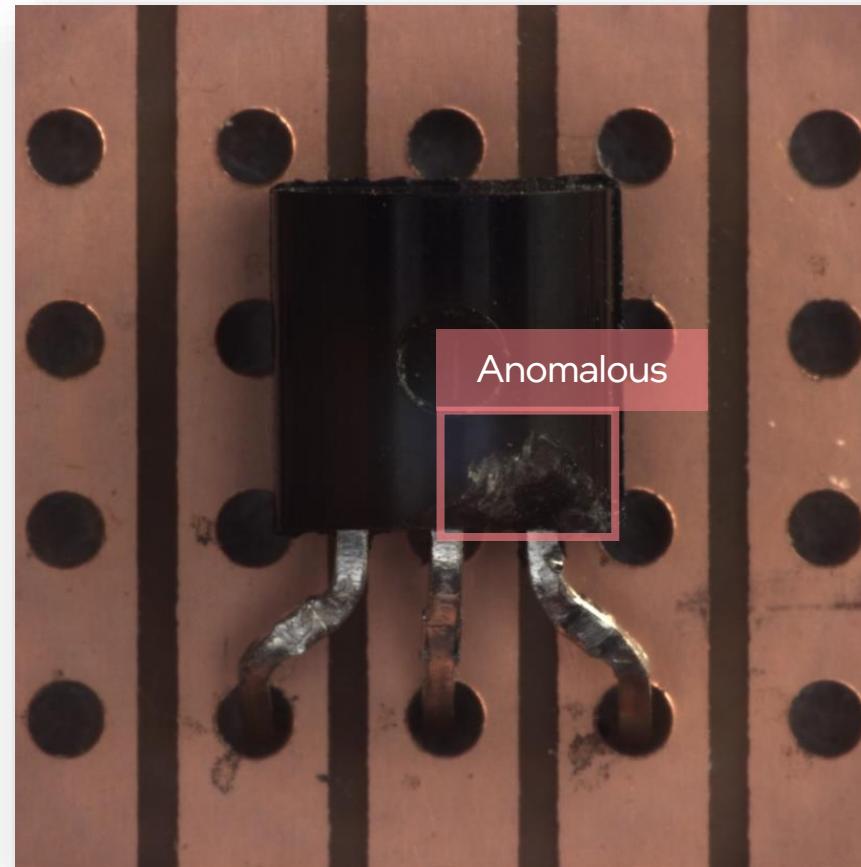
W. Liu and W. Luo, D. Lian and S. Gao. "Future Frame Prediction for Anomaly Detection -- A New Baseline." CVPR 2018.
Lu, C., Jianping S., and Jiaya J., Abnormal event detection at 150 fps in matlab, ICCV 2013

Supported Tasks

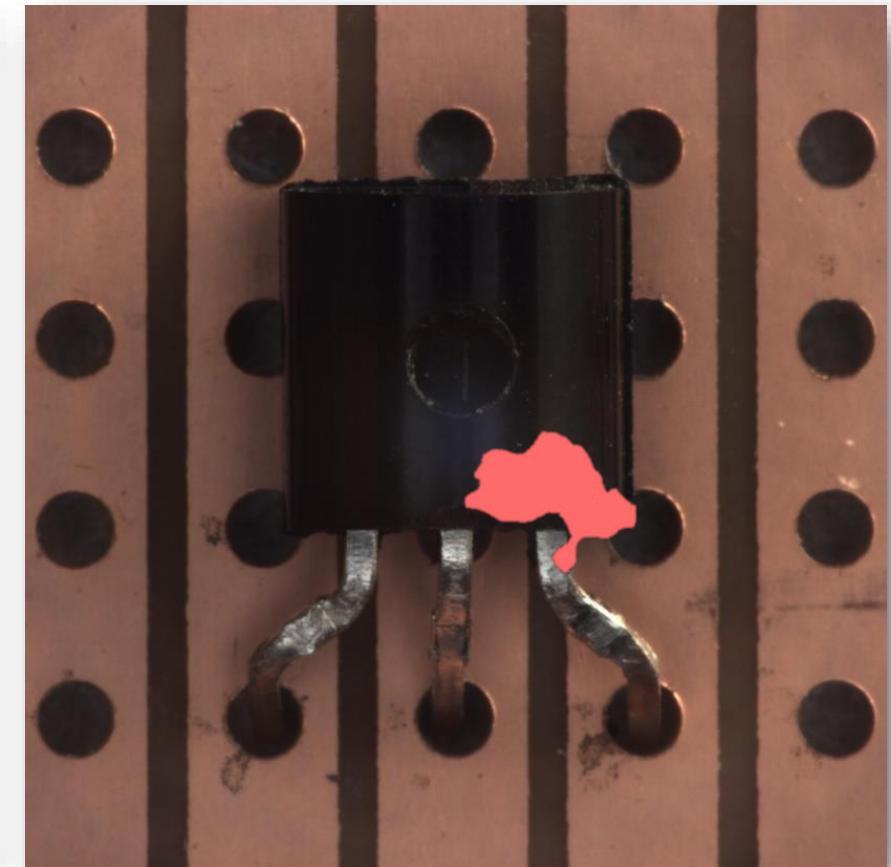
Anomalous



Classification

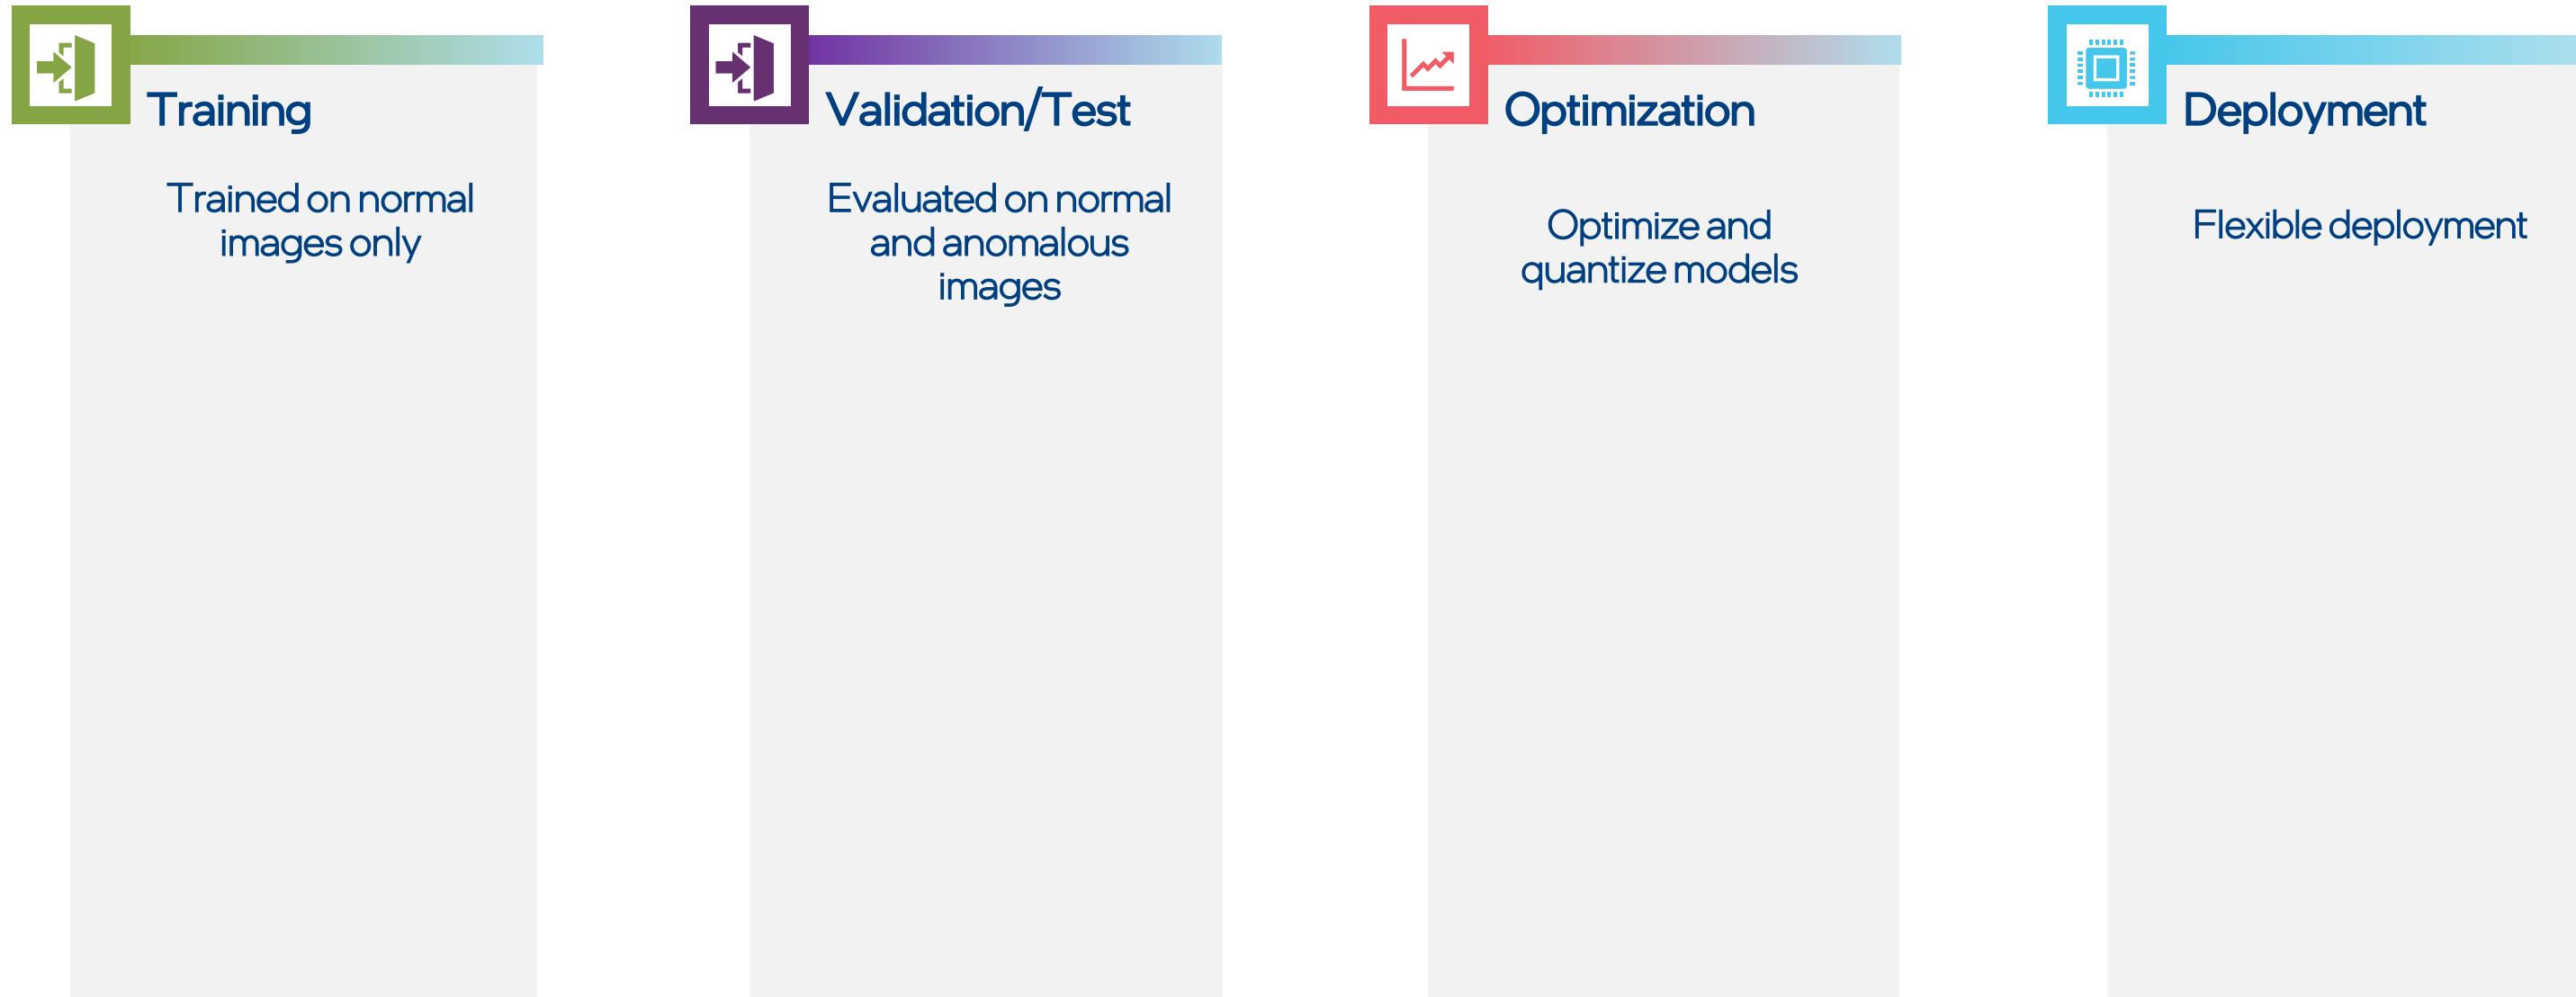


Detection

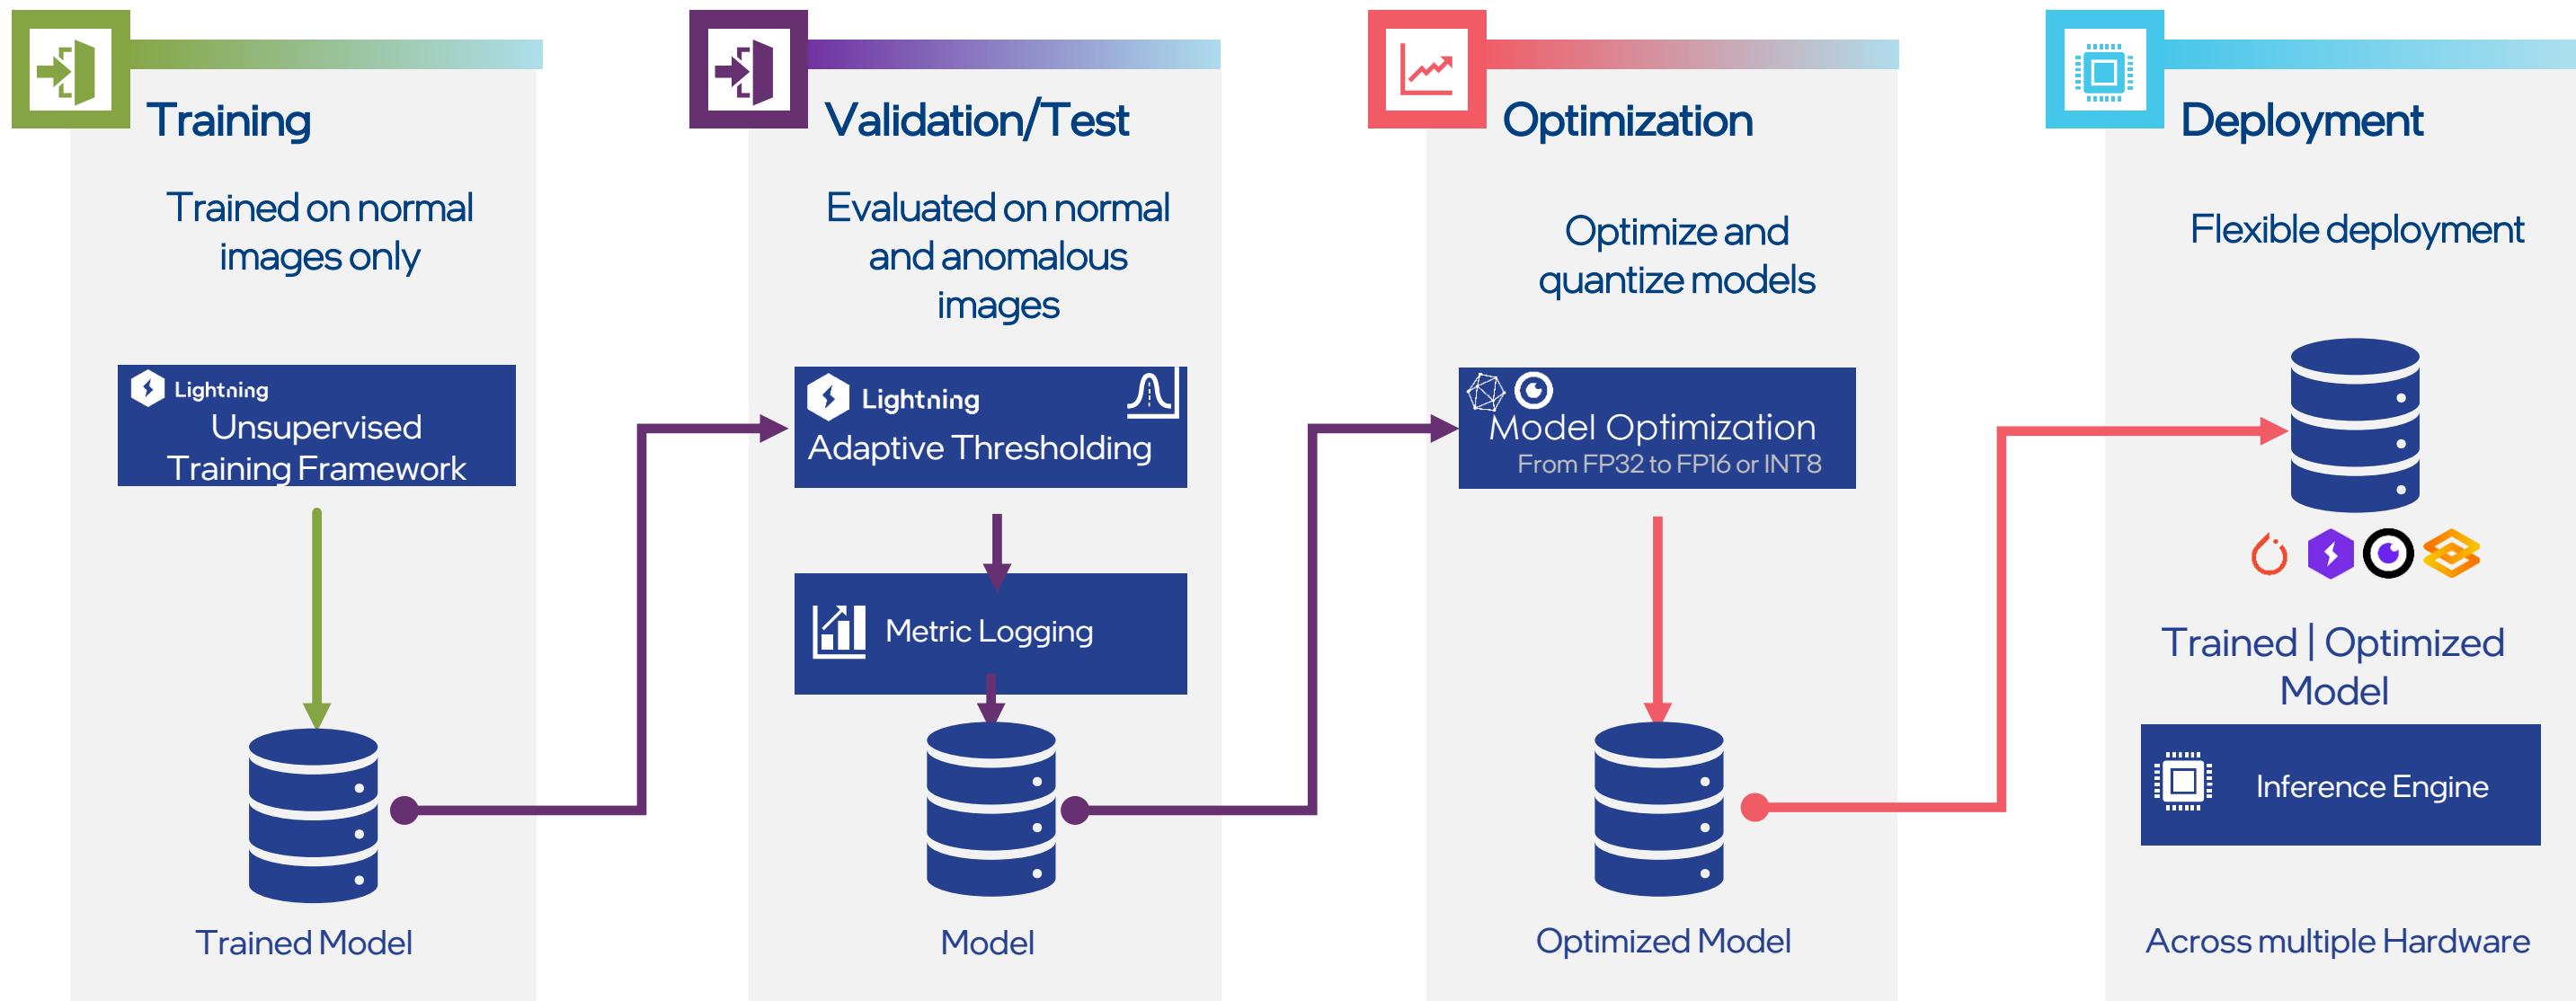


Segmentation

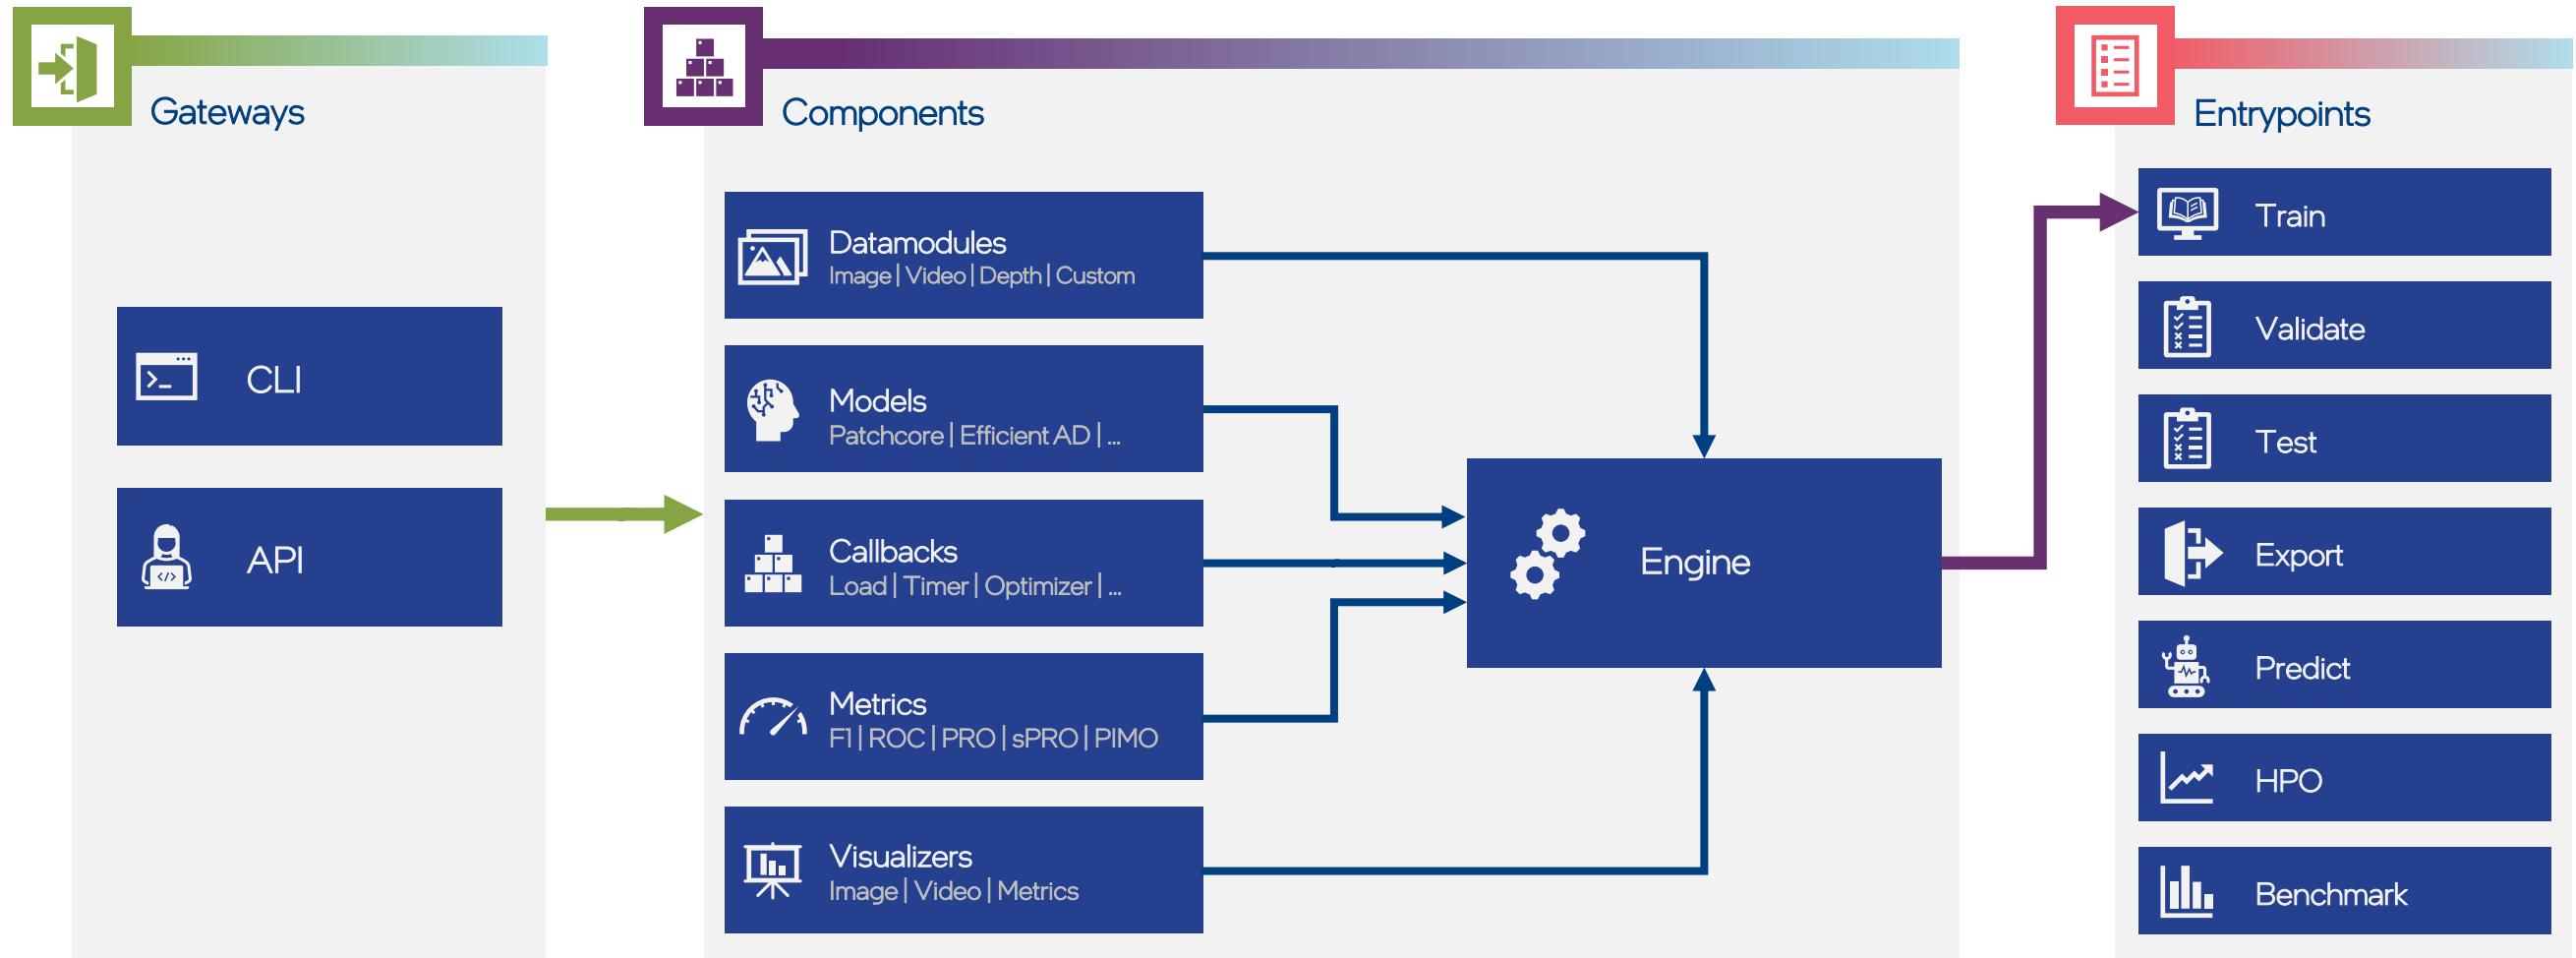
High-Level Workflow



High-Level Workflow



High-Level Architecture



Entry points - Brief



API – Entrypoints

● ● ●

```
# Train the model  
>>> engine.train(...)  
  
# Validate the model  
>>> engine.validate(...)  
  
# Test the model  
>>> engine.test(...)  
  
# Export the model  
>>> engine.export(...)  
  
# Make a prediction  
>>> engine.predict(...)
```



CLI – Entrypoints

● ● ●

```
# Train the model  
> anomilib train ...  
  
# Validate the model  
> anomilib validate ...  
  
# Test the model  
> anomilib test ...  
  
# Export the model  
> anomilib export ...  
  
# Make a prediction  
> anomilib predict ...
```

Learn and build step-by-step

01 Installation

02 Data Acquisition

03 Training

04 Testing

05 Optimization

06 Deployment

01 Installation



Easy Installation

Hardware-agnostic, configurable installation

• • •

```
> pip install anomalib
> anomalib install -h
<- Arguments _____
  Usage: anomalib [options] install [-h]
        [--option {full,core,dev,logge
        [-v]

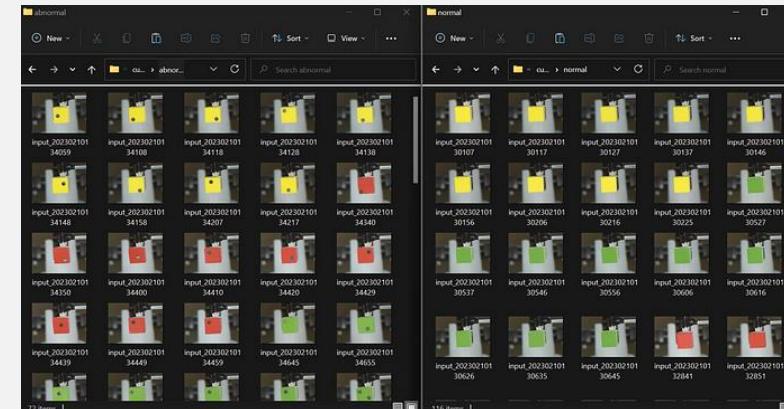
  Options:
    -h, --help Show this help message and exit.
    --option {full,core,dev,loggers,notebooks,openvino}
      Install the full or optional-dependencies
      (type: None, default: full)
    -v, --verbose Set Logger level to INFO (default: False)
```

02 Data Acquisition

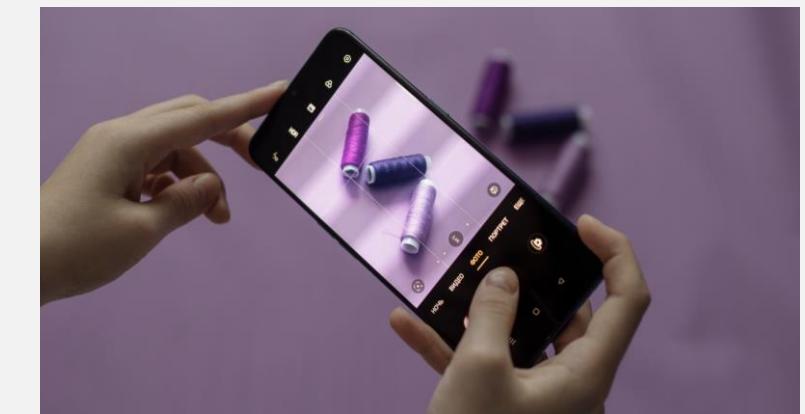
Option 1:
Create a dataset using
Dobot Magician



Option 2:
Download images
[-Our dataset](#)
[-MVTec](#)



Option 3:
Create a dataset with your
own phone or camera



02 Data Acquisition



Datamodule API

Public benchmark support

```
>>> from anomalib.data import MVTec, BTech, Visa,  
...  
>>> datamodule = MVTec()  
>>> datamodule.setup()  
  
>>> i, train_data =  
next(enumerate(datamodule.train_dataloader()))  
>>> print(train_data.keys())  
dict_keys(['image_path', 'label', 'image'])  
  
>>> print(train_data["image"].shape)  
torch.Size([32, 3, 256, 256])  
  
>>> i, val_data =  
next(enumerate(datamodule.val_dataloader()))  
>>> i, test_data =  
next(enumerate(datamodule.test_dataloader()))
```

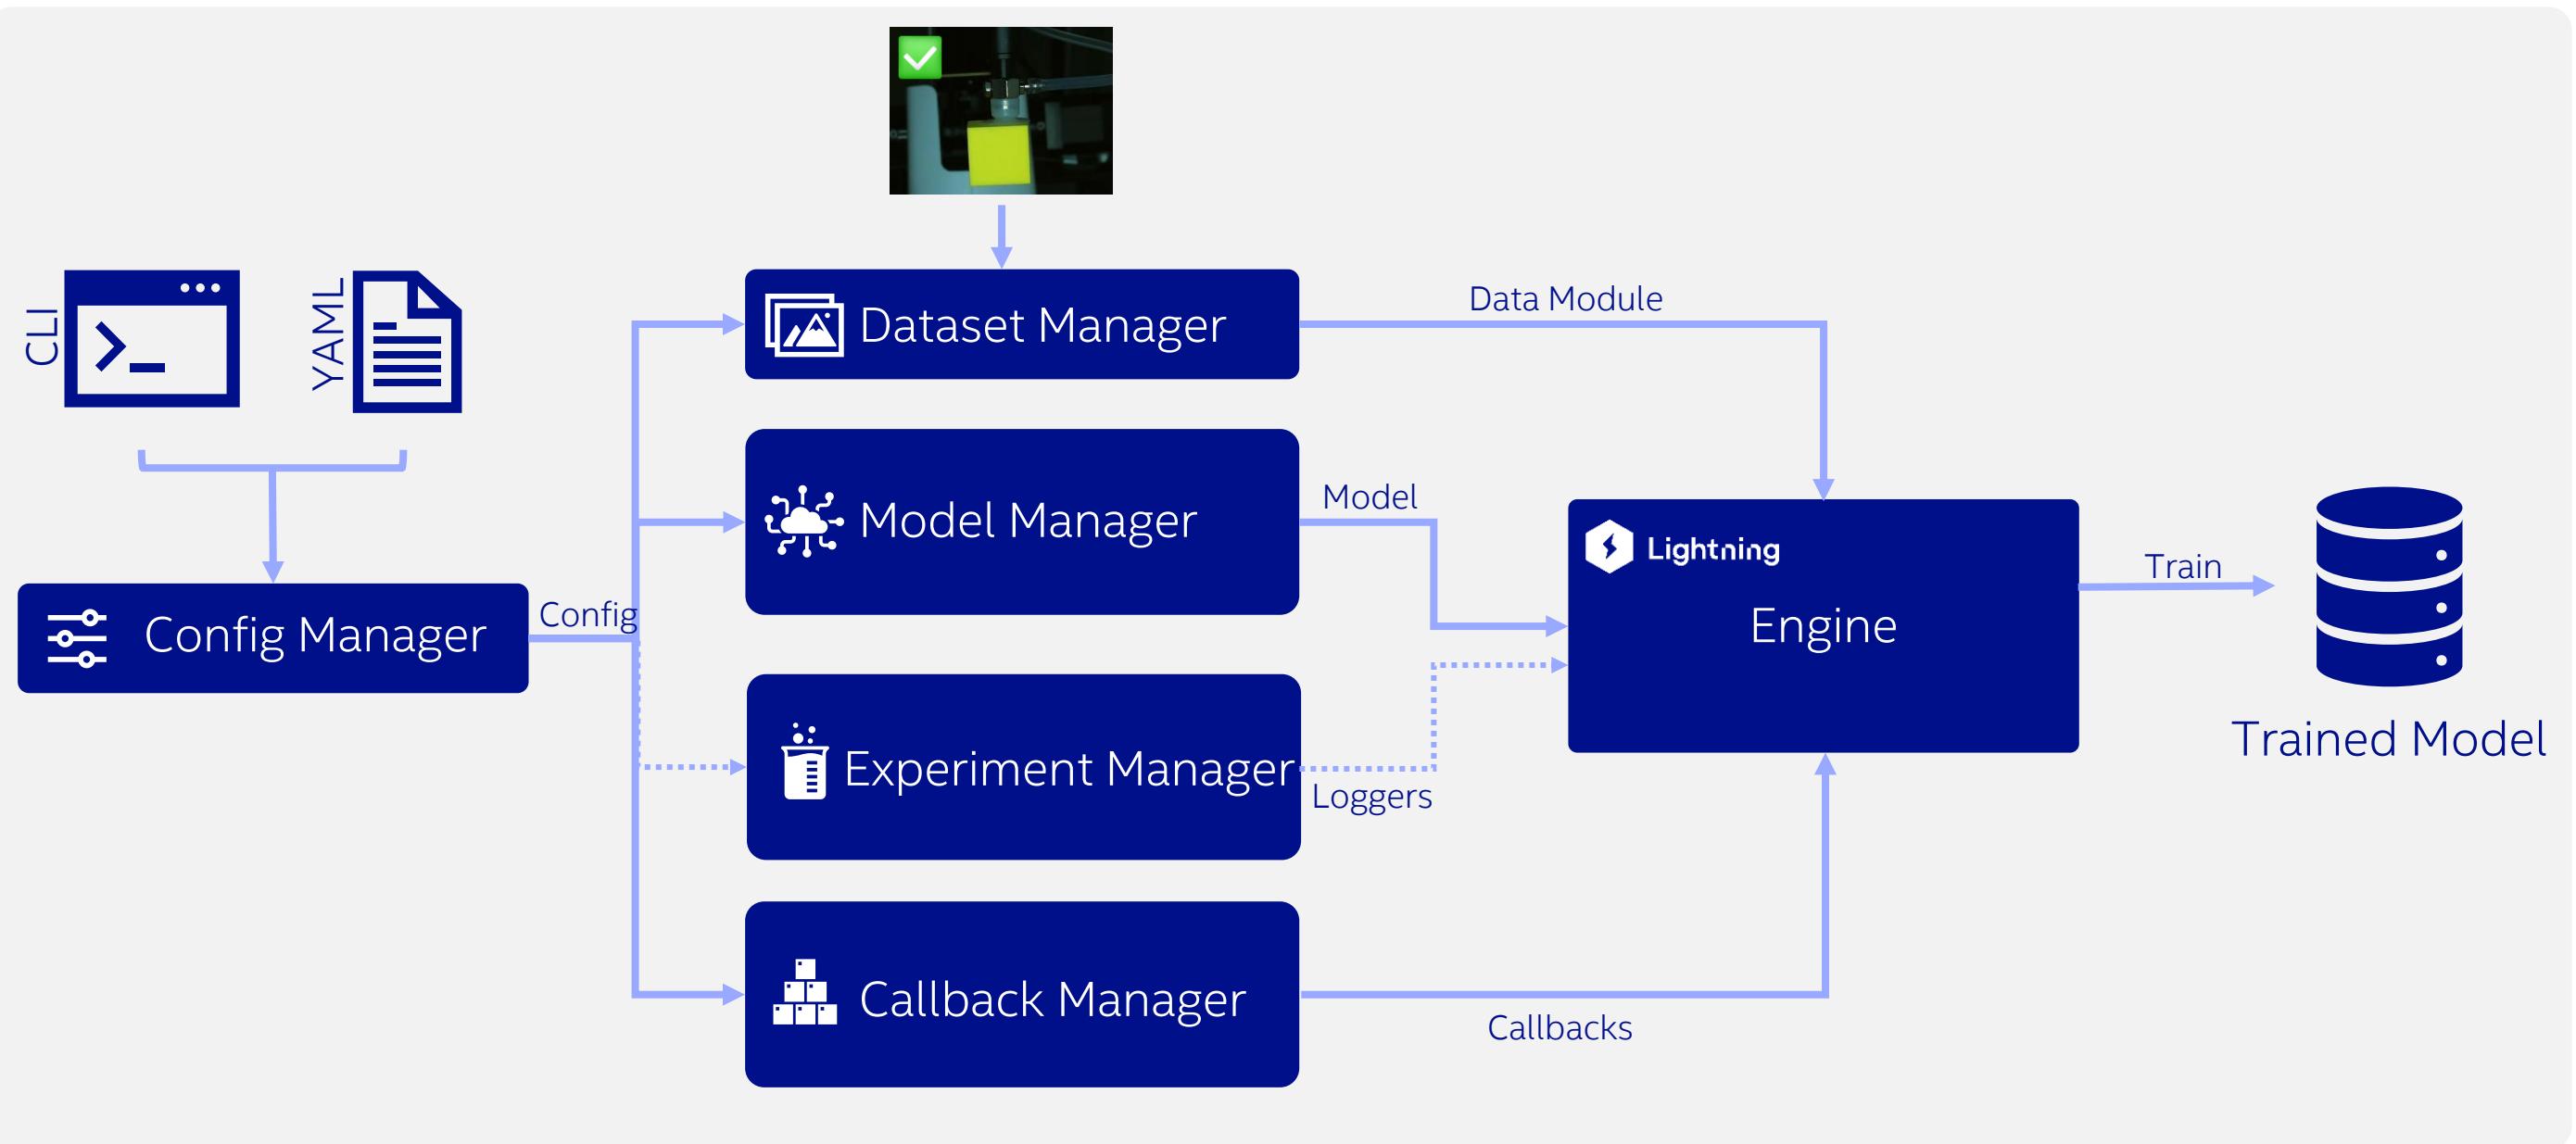


Datamodule API

Custom data support

```
>>> from anomalib.data import Folder  
>>> datamodule = Folder(  
...     name="my_custom_dataset",  
...     root="datasets/hazelnut",  
...     normal_dir="good",  
...     abnormal_dir="bad",  
...     task="classification"  
... )  
>>> datamodule.setup()  
  
>>> i, train_data =  
next(enumerate(datamodule.train_dataloader()))  
  
>>> print(train_data["image"].shape)  
torch.Size([32, 3, 256, 256])
```

03 Training



03 Training



* This is a growing list as anomalib is constantly updated with newer models

Entry points



API - Train

● ● ●

```
# Import the required modules
>>> from anomalib.data import MVTec
>>> from anomalib.models import Patchcore
>>> from anomalib.engine import Engine

# Initialize the datamodule, model and engine
>>> datamodule = MVTec()
>>> model = Patchcore()
>>> engine = Engine()

# Train the model
>>> engine.train(datamodule=datamodule, model=model)
```



CLI - Train

● ● ●

```
# Get help about the training arguments, run:
> anomalib train -h

# Train by using the default values.
> anomalib train \
    --data anomalib.data.MVTec \
    --model Patchcore

# Train by overriding arguments.
> anomalib train \
    --model Patchcore \
    --data anomalib.data.MVTec \
    --data.category transistor

# Train using YAML config file.
> anomalib train \
    --config patchcore.yaml
```

Metrics



Metrics API

Image and Pixel-wise metric support via torchmetrics

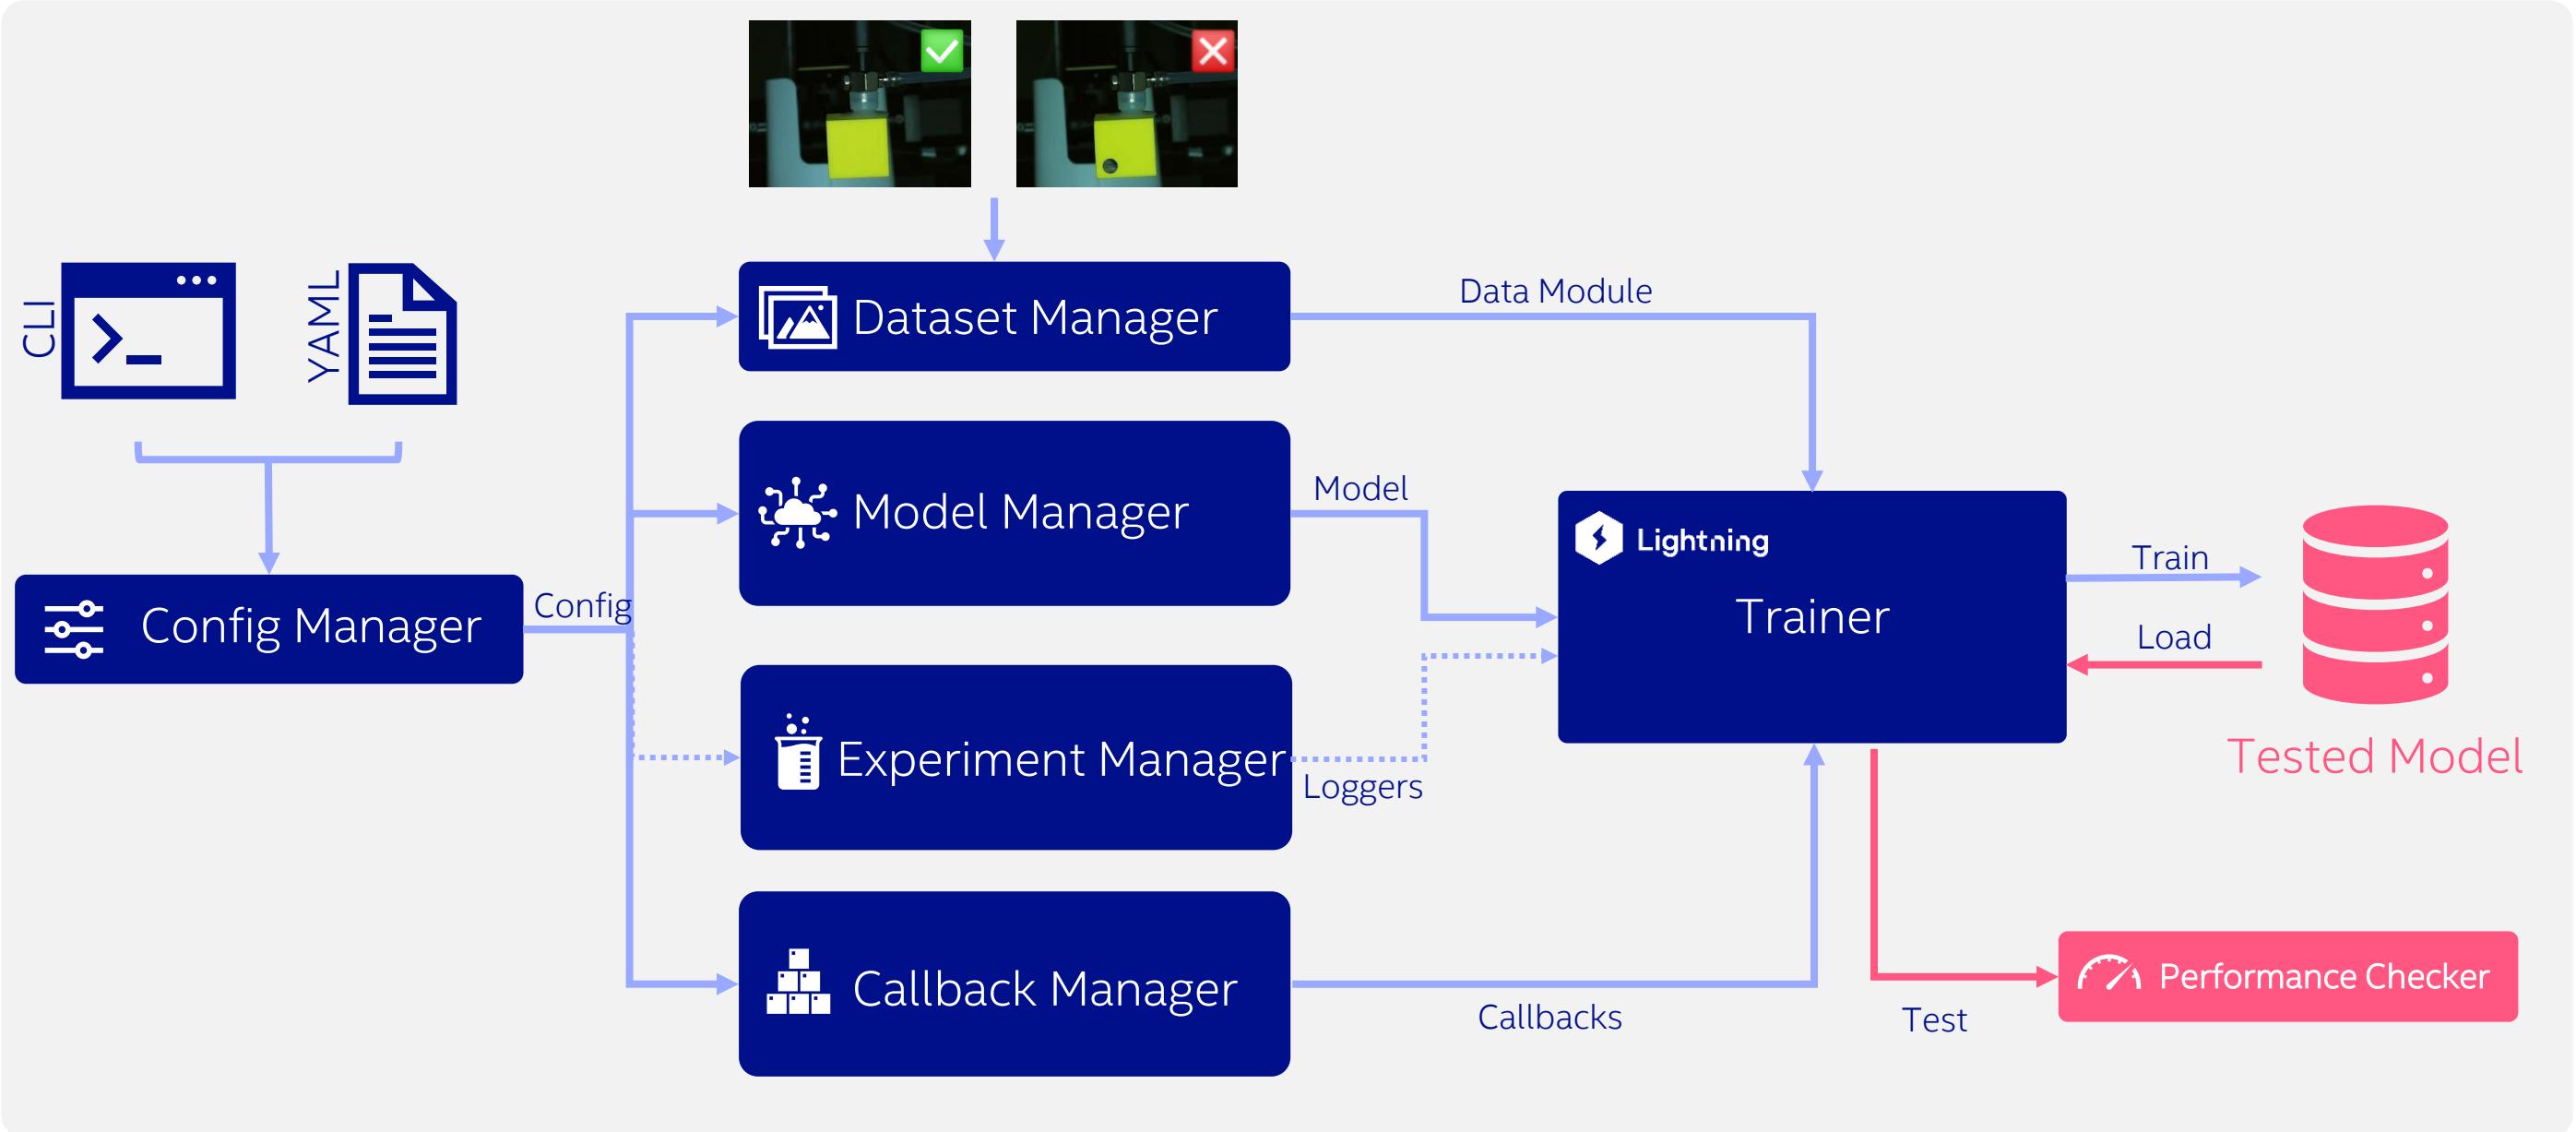


```
# Import the required modules
>>> from anomalib.data import MVTec
>>> from anomalib.models import Patchcore
>>> from anomalib.engine import Engine

# Customize the image and/or pixel metrics
>>> datamodule = MVTec()
>>> model = Patchcore()
>>> engine = Engine(pixel_metrics=["F1Score", "AUROC", "AUPRO"])

# Train the model
>>> engine.train(datamodule=datamodule, model=model)
```

04 Testing



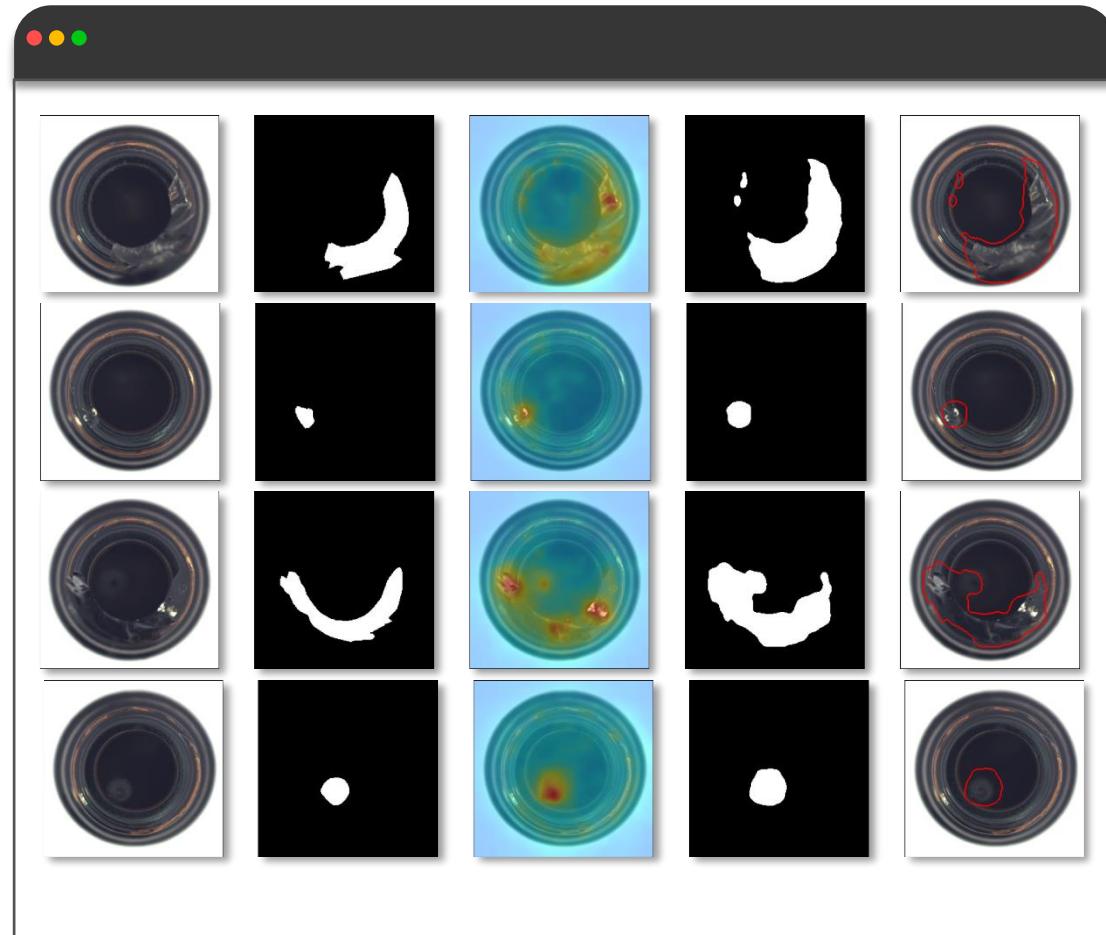
03 Training

Modular API

```
# Import data, model and trainer
from anomalib.data import MVTec
from anomalib.models import Padim
from anomalib.trainer import AnomalibTrainer

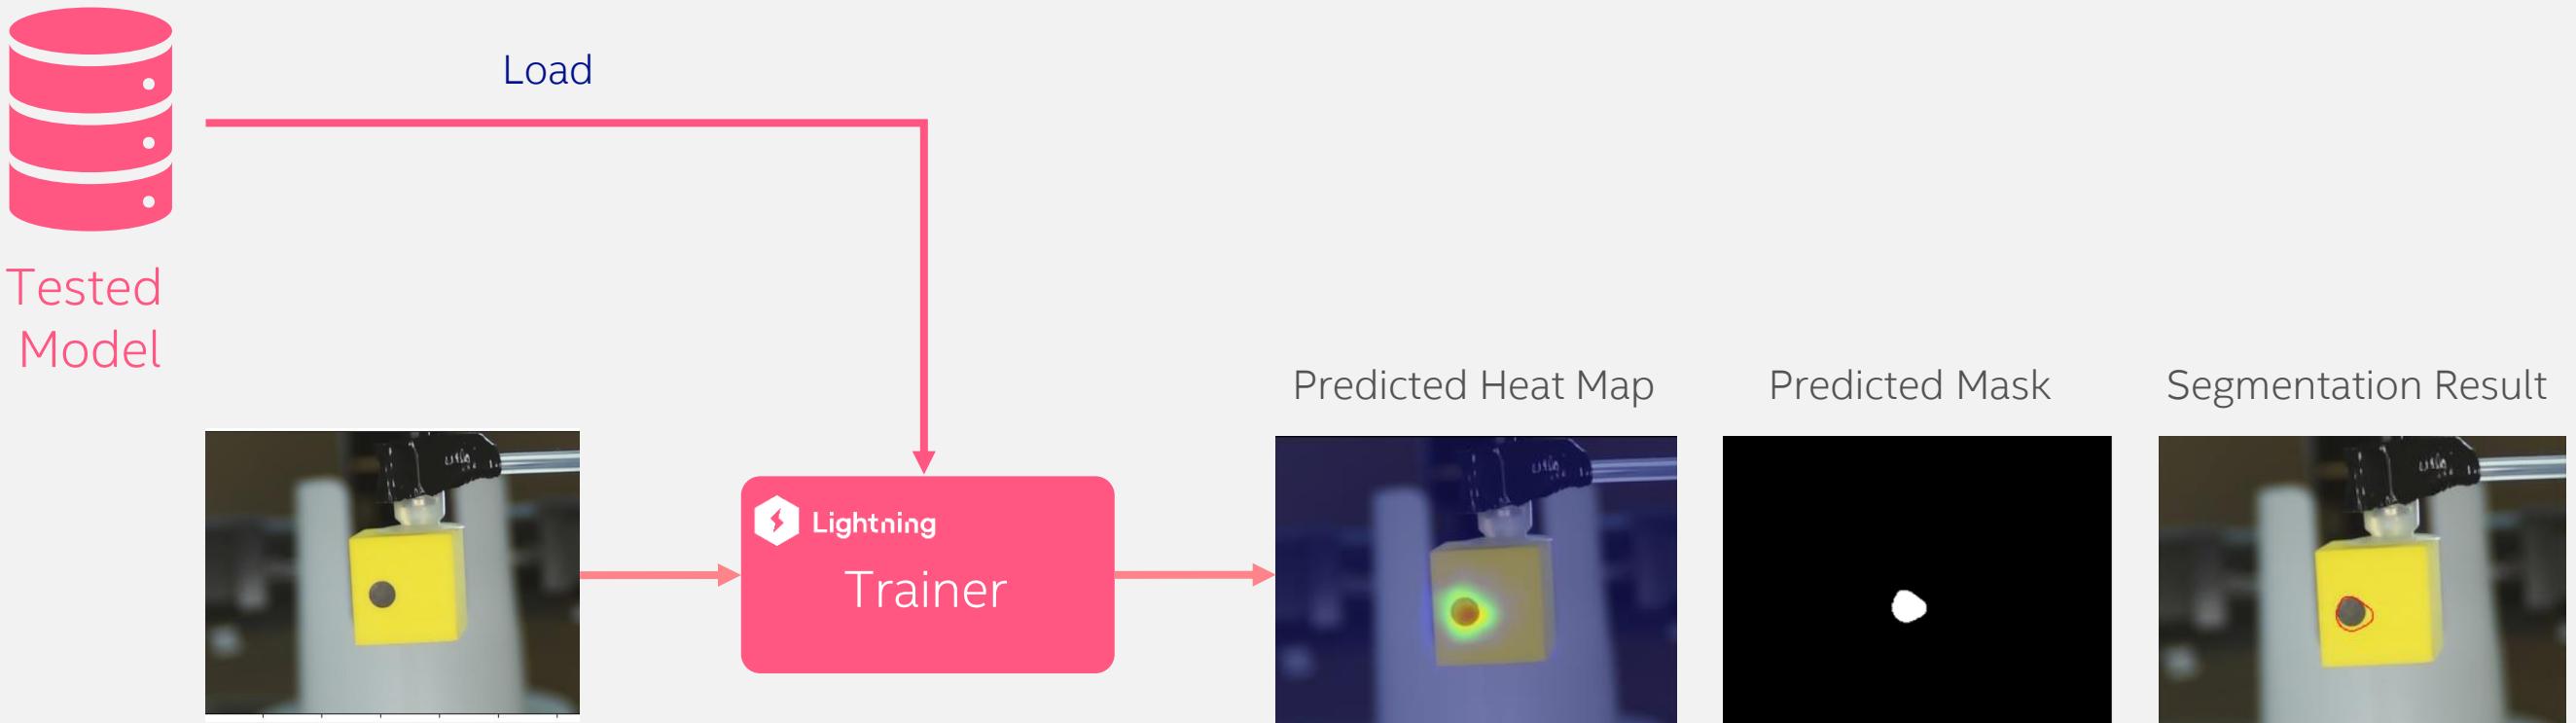
# Initialize the data, model and trainer
datamodule = MVTec("./MVTec/", "bottle", image_size=256)
model = Padim(input_size=(256, 256))
trainer = AnomalibTrainer(pixel_metrics=["AUROC"], max_epochs=1)

# Train the model
trainer.fit(model, datamodule=datamodule)
~/Pr/anomalib feature/custom_loops !1 ?2
```



* This is a growing list as anomalib is constantly updated with newer models

04 Testing



Entry points



API – Test / Export

● ● ●

```
# Training
...
# Test the model
>>> engine.test(datamodule=datamodule, model=model)
# Export the model
>>> model_path = engine.export(
...     model=model, # model from the previous step
...     export_type="openvino", # "torch", "onnx", "openvino"
... )
```



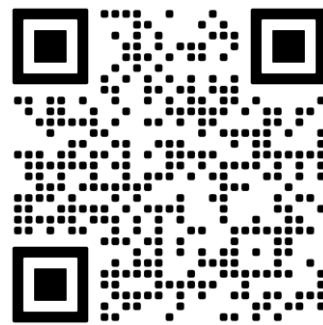
CLI – Test / Export

● ● ●

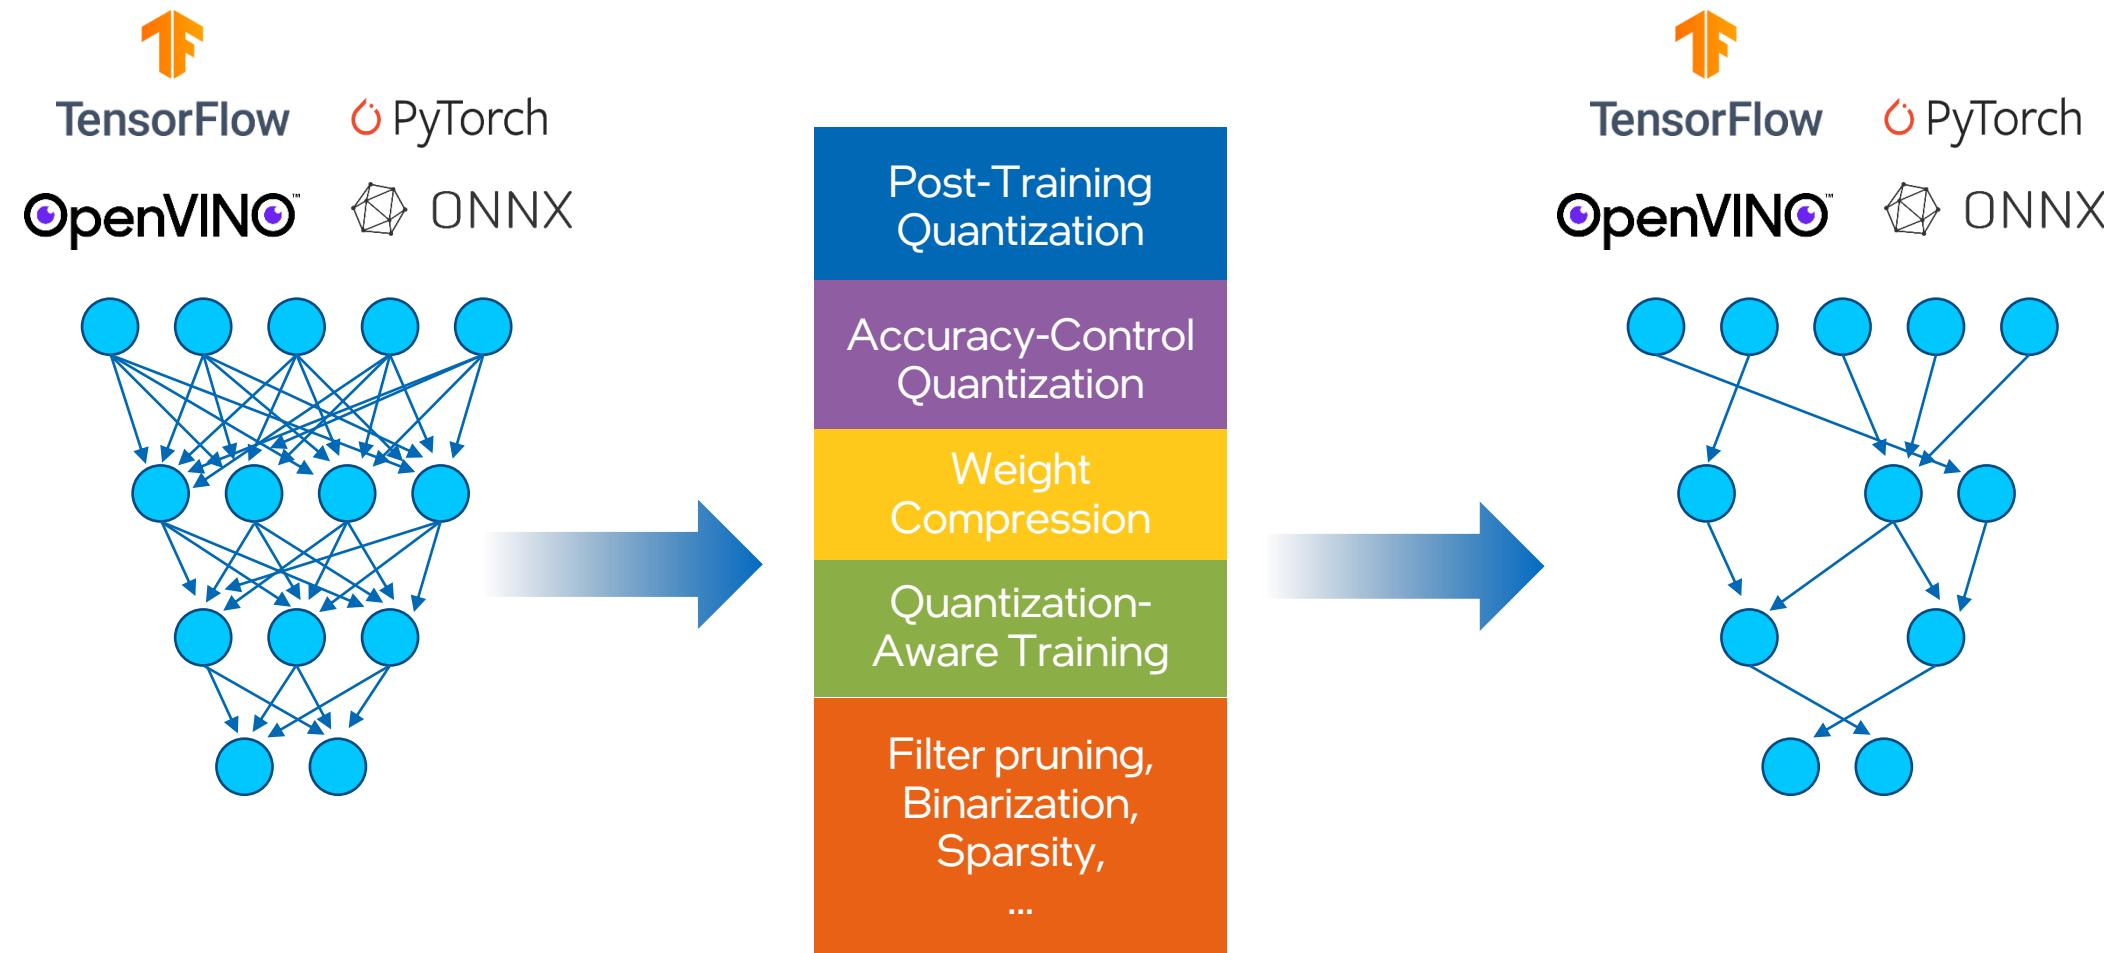
```
# Get help about the training arguments, run:
> anomalib test -h

# Train by using the default values.
> anomalib test \
    --data anomalib.data.MVTec \
    --model Patchcore

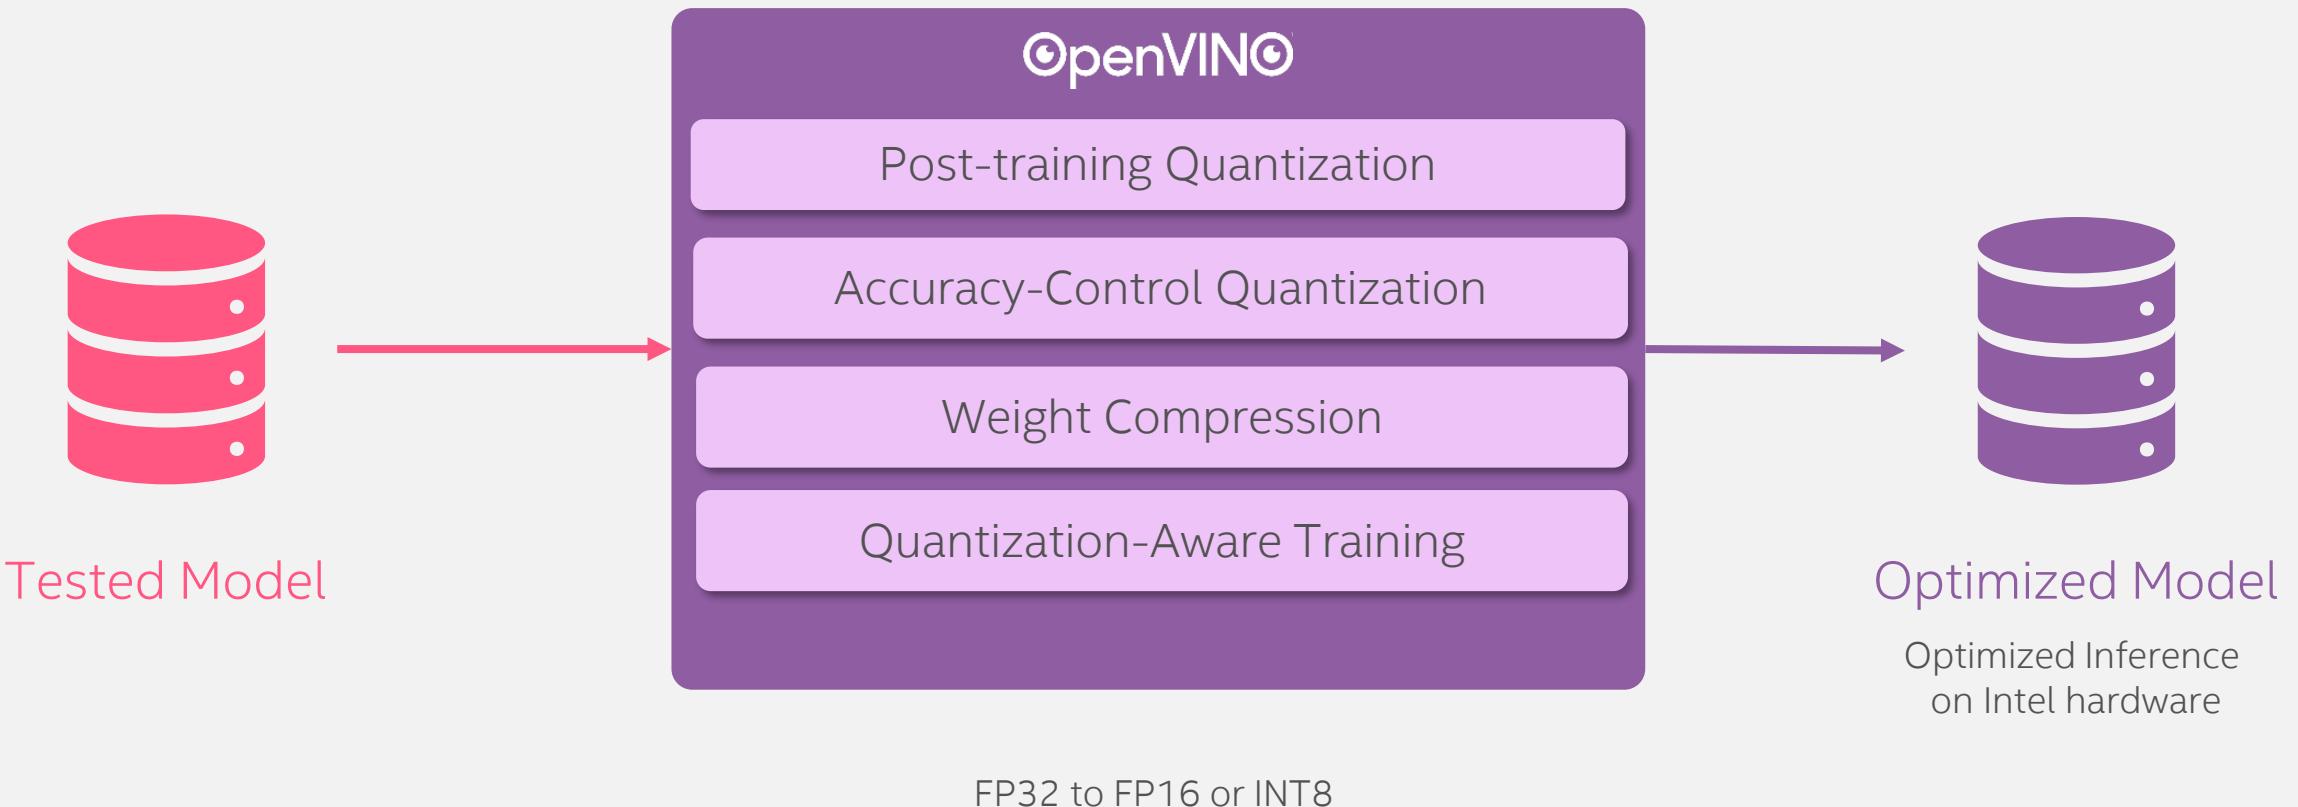
# Train by overriding arguments.
> anomalib export \
    --model Patchcore \
    --data anomalib.data.MVTec \
    --export_type openvino \
```

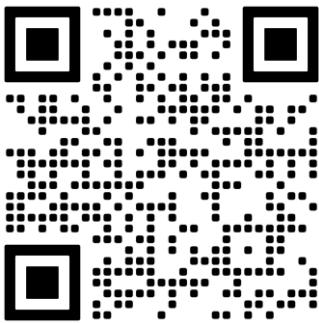


Neural Network Compression Framework (NNCF)



05 Optimization





Neural Network Compression Framework (NNCF)



API – Model Compression and quantization

● ● ●

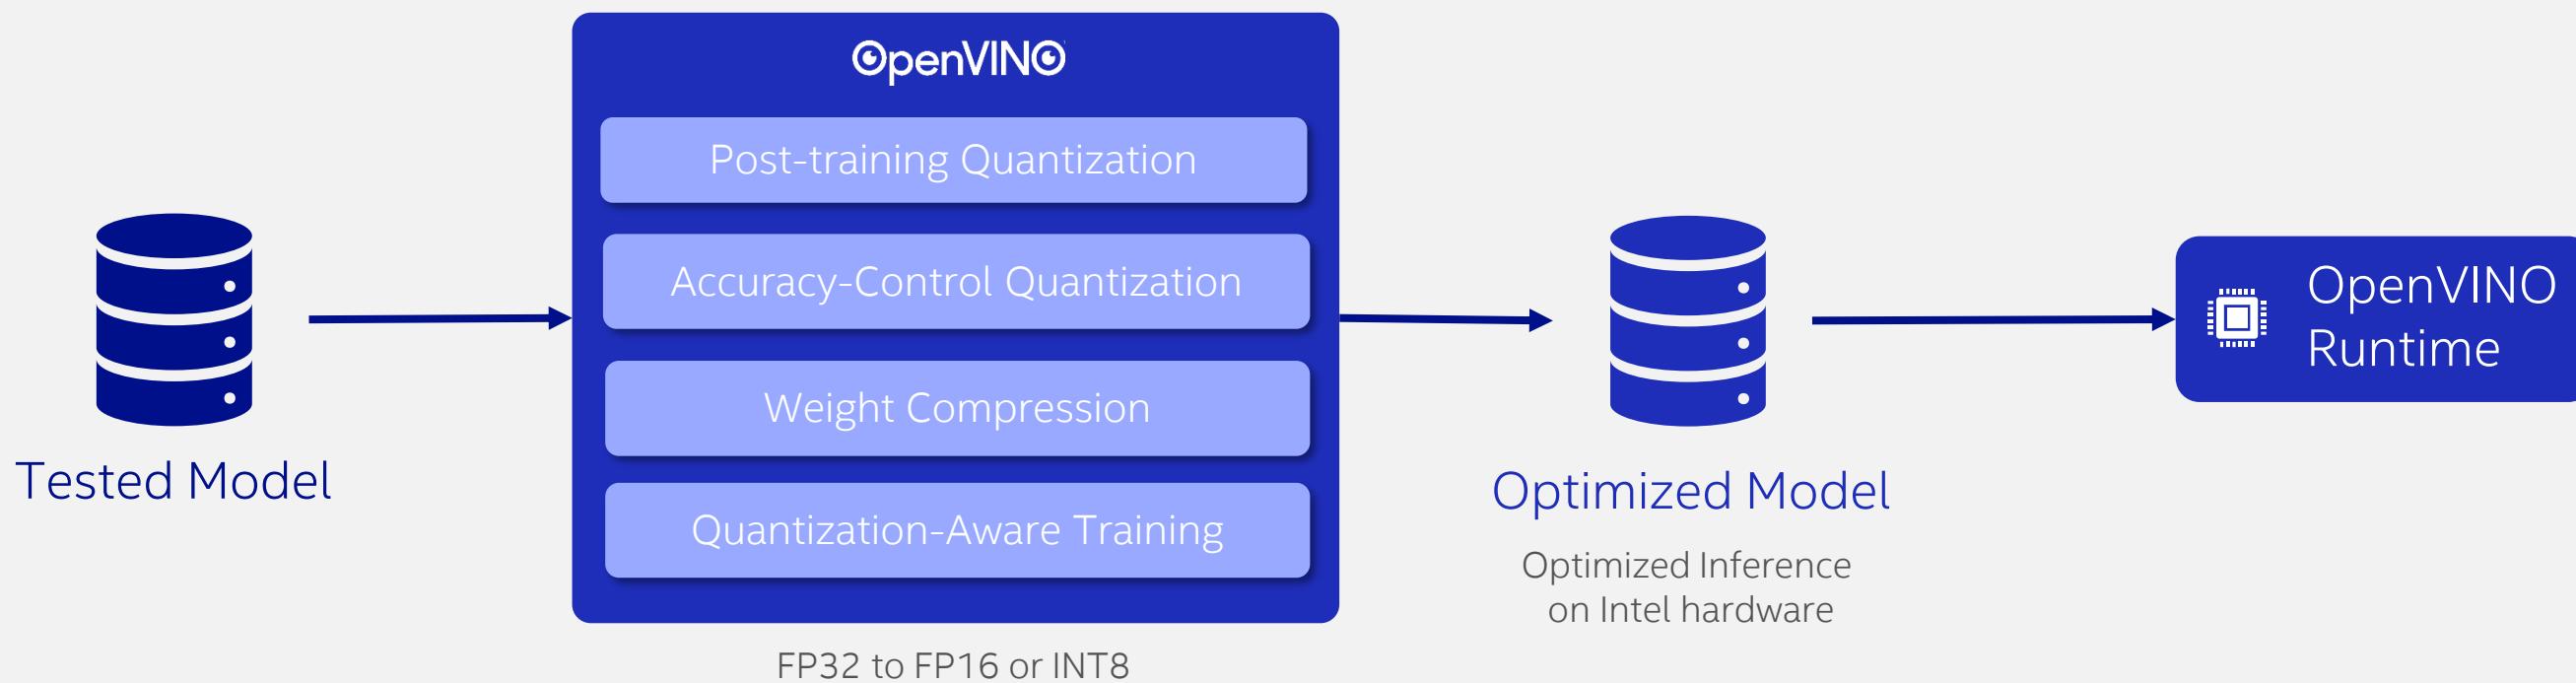
```
# Post-Training Quantization
>>> engine.export(model, ExportType.OPENVINO, export_dir,
compression_type=CompressionType.INT8_PTQ, datamodule=datamodule)

# Accuracy-Control Quantization
>>> engine.export(model, ExportType.OPENVINO, export_dir,
compression_type=CompressionType.INT8_ACQ, datamodule=datamodule),
metric="F1Score")

# Weight Compression
>>> engine.export((model, ExportType.OPENVINO, export_dir,
compression_type=CompressionType.FP16, datamodule=datamodule)

>>> engine.export((model, ExportType.OPENVINO, export_dir,
compression_type=CompressionType.INT8, datamodule=datamodule)
```

06 Deployment

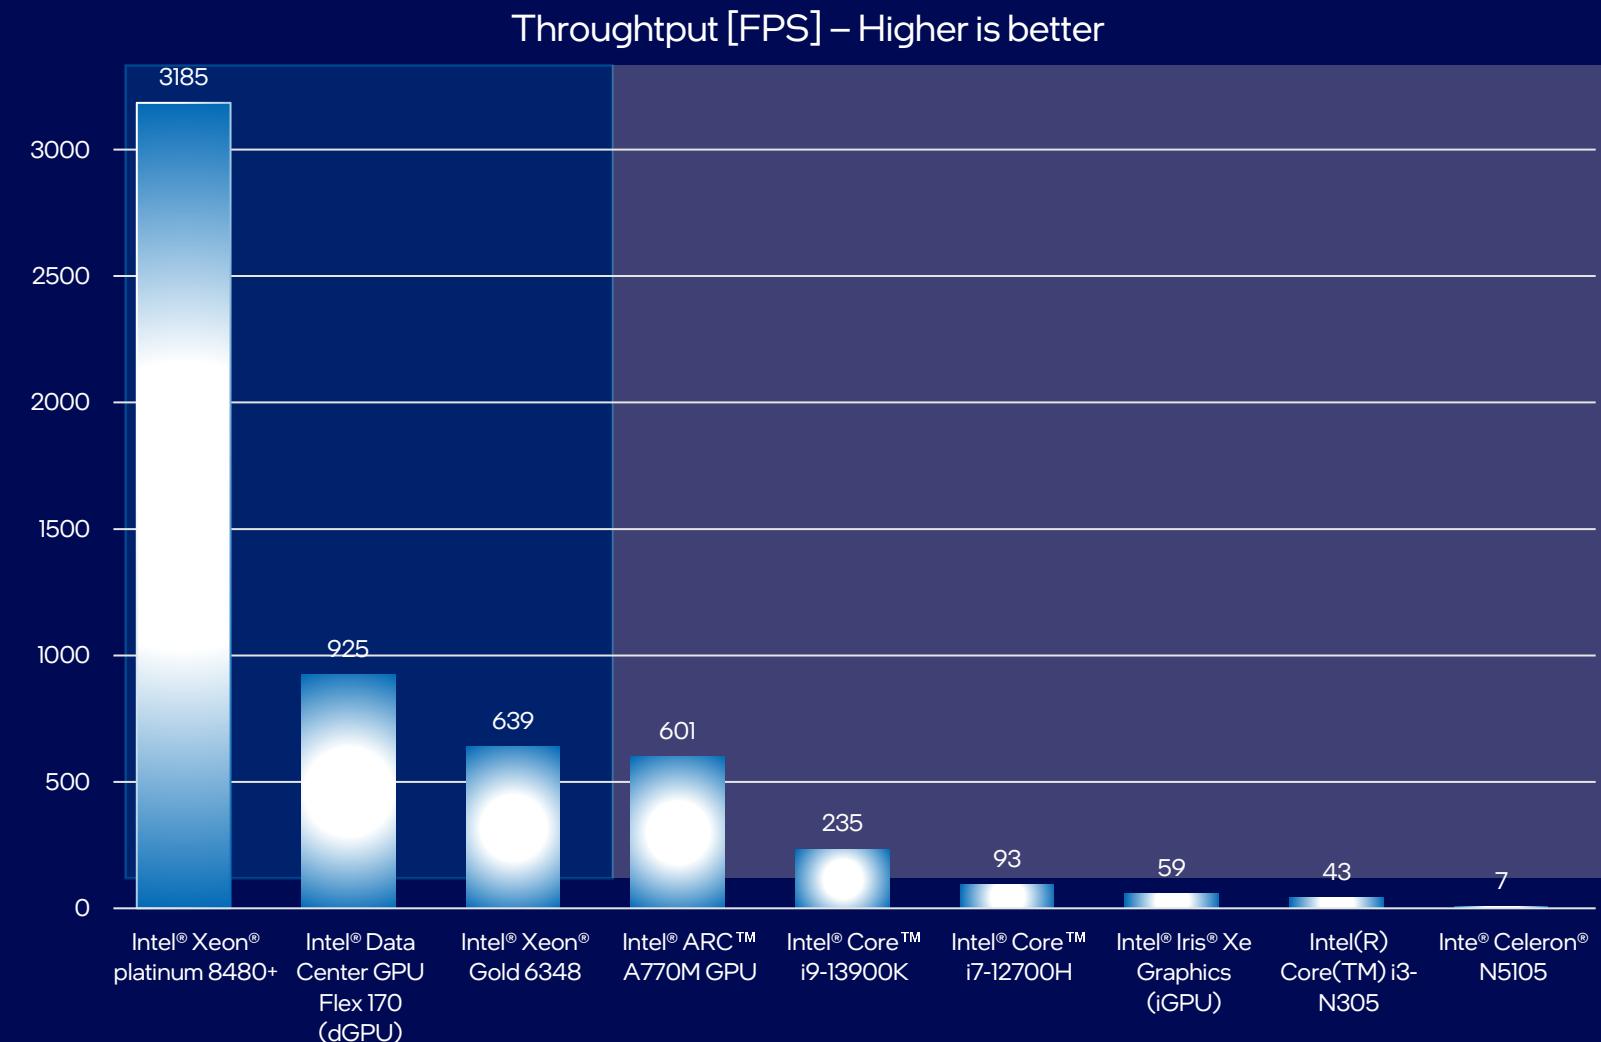




Demo using Anomalib Notebooks



Optimized Performance



Inferencing results on Padim model
Trained using Anomalib

Results may vary

Platform Configurations for Performance Benchmarks

System Board	Intel Corporation D50DNP1SBB	Intel Corporation M50CYP2SBSTD	Intel® Client Systems NUC13RNGi9	Intel® Client Systems NUC12SNKi72	AAEON UPN-ADLN01 V1.0 220950173	AIXBoard Aixboard-jspl
CPU	Intel(R) Xeon(R) Platinum 8480+	Intel(R) Xeon(R) Gold 6348 CPU @ 2.60GHz	13th Gen Intel(R) Core(TM) i9-13900K @ 5.4 GHz	12 th Gen Intel® Core™ i7-12700H @ 2.30 GHz	Intel® Core™ i3-N305 @ 3.80 GHz	Intel® Celeron® N5105 @ 2.00GHz
Sockets/Physical cores	1 / 56 (112 Threads)	2 / 28 (56 Threads)	1 / 24 (32 Threads)	1 / 14 (20 Threads)	1 / 8 (8 Threads)	1 / 4
HyperThreading/Turbo Setting	Enabled / On	Enabled / On	Enabled / On	Enabled / On	Disabled	Disabled
Memory	512 GB DDR4 @ 4800 MHz	256 GB DDR4 @ 3200 MHz	64 GB DDR5 @ 4800 MHz	64 GB DDR4 @ 3200 MHz	16GB DDR5 @4800 MHz	8GB LPDDR4 @ 2933MHz
OS	Ubuntu 22.04.2 LTS	Ubuntu 22.04.2 LTS	Ubuntu 22.04.2 LTS	Windows 11 Enterprise v22H2	Ubuntu 22.04.2 LTS	Ubuntu 20.04.6 LTS
Kernel	5.15.0-72-generic	5.15.0-57-generic	5.15.0-73-generic	22621.1702	5.15.0-1028-intel-iotg	5.15.0-1016-intel-iotg
Software	OpenVINO 2023.0.0	OpenVINO 2022.3.0	OpenVINO 2022.3.0	OpenVINO 2022.3.0	OpenVINO 2022.3	OpenVINO 2022.3
BIOS	Intel Corp. SE5C7411.86B.9525.D13.2302 071333	Intel Corp. SE5C620.86B.01.01.0007.221 0270543	Intel Corp. SBRPL579.0053.2022.1125.01 01	Intel Corp. SNADL357.0053.2022.1102.12 18	American Megatrends International, LLC. UNADAM10	AIXBoard Aix-jspl-v1
BIOS Release Date	02/07/2023	10/27/2022	11/25/2022	11/02/2022	12/15/2022	03/06/2023
GPU	N/A	1x Intel® Data Center GPU Flex 170	1x Intel® Arc A770 16GB, 512 EU 1x Intel® UHD Graphics 770	1x Intel® Arc A770 ™ 16GB, 512 EU 1x Intel® Iris® Xe Graphics	N/A	N/A
Workload: Codec, resolution, frame rate Model, size (HxW), BS	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32 (CPU) FP16 (GPU)	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32 (CPU)	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32 (CPU) FP16 (GPU)	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32	Padim Model – Backbone ResNet-18, input size [256, 256] , batch 32 FP32
TDP	350W	235W	150W	45W	15W	10W
Benchmark Date	May 31, 2023	May 29, 2023	May 29, 2023	May 29, 2023	May 29, 2023	May 29, 2023
Benchmarked by	Intel Corporation	Intel Corporation	Intel Corporation	Intel Corporation	Intel Corporation	Intel Corporation

VAND 2.0 Challenge at CVPR

Sponsored by Intel

The Challenge

This year our challenge aims to bring visual anomaly detection closer to industrial visual inspection, which has wide real-world applications. This challenge consists of two categories:

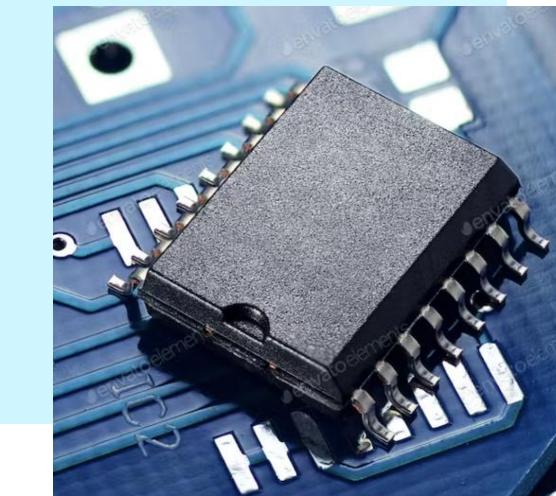
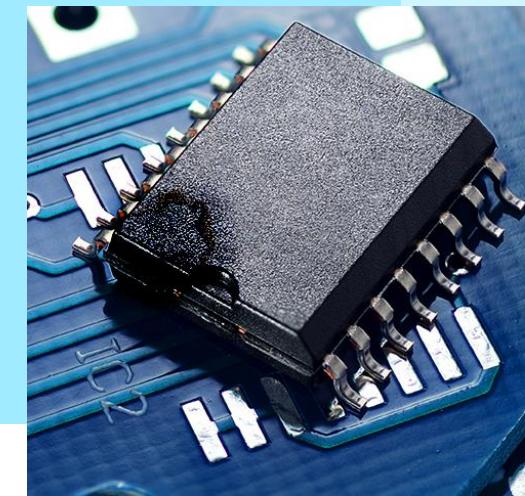
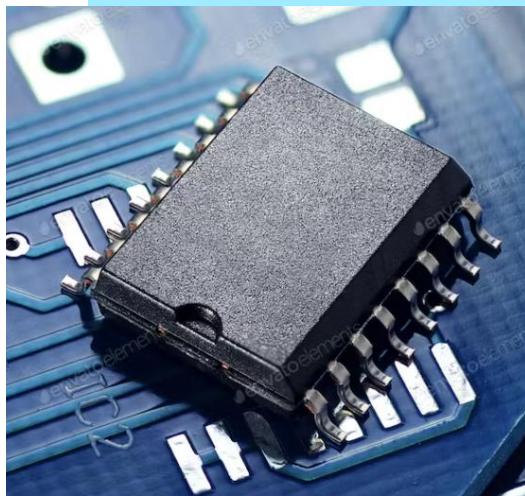
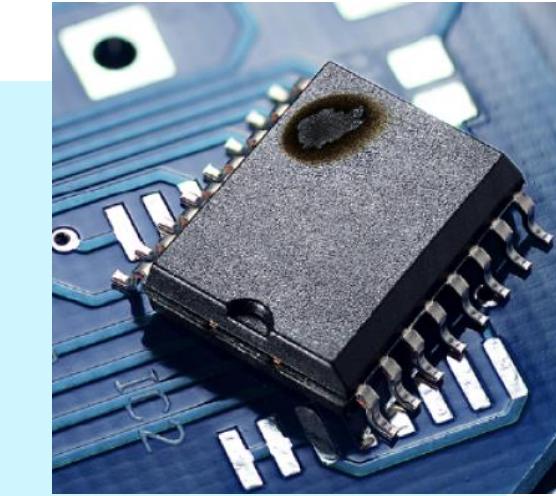
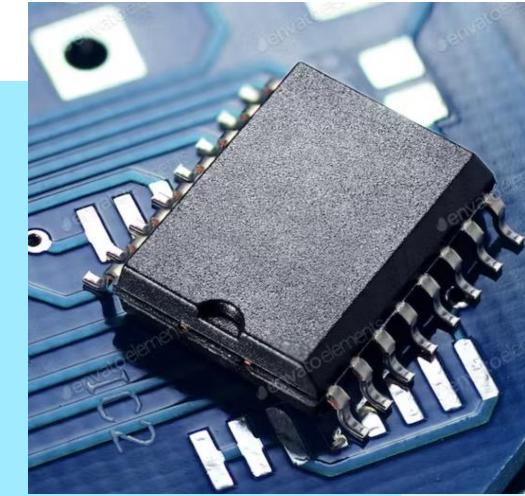
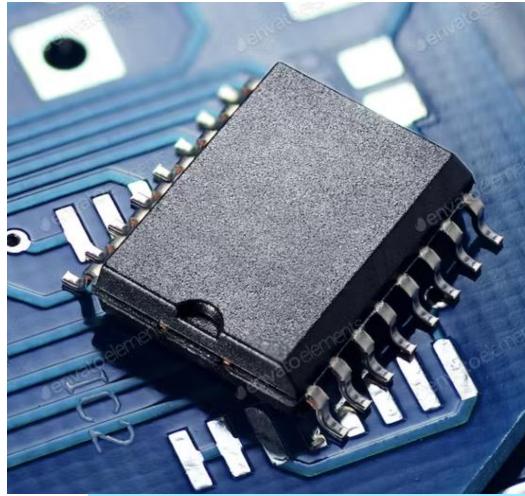
- **Category 1**— Adapt & Detect: Robust Anomaly Detection in Real-World Applications
- **Category 2**— VLM Anomaly Challenge: Few-Shot Learning for Logical and Structural Detection



420 participants 62 submissions



What / Where are the defects? – Category 1



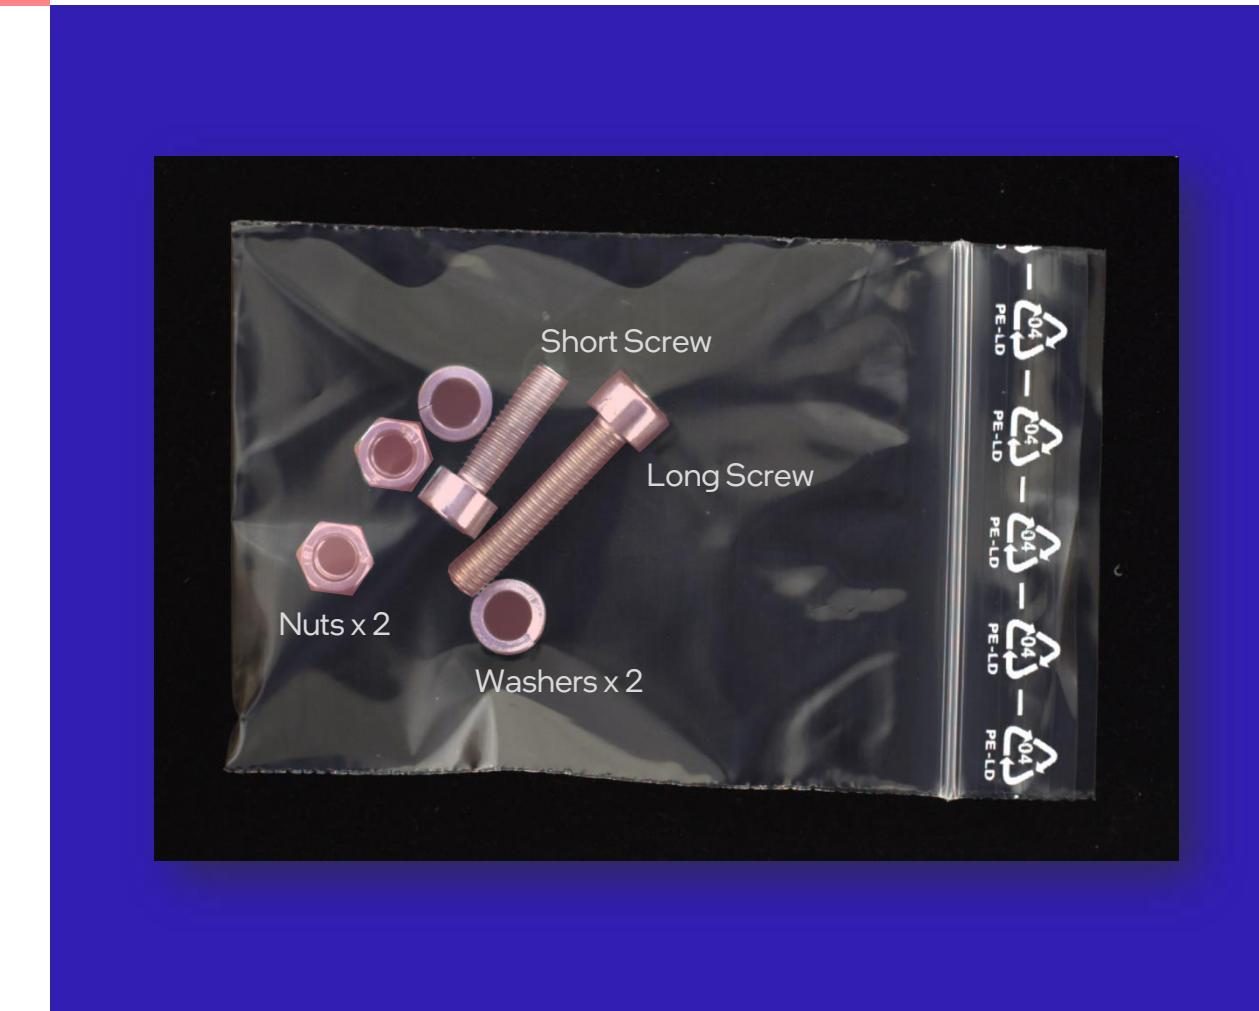
What / Where are the defects?



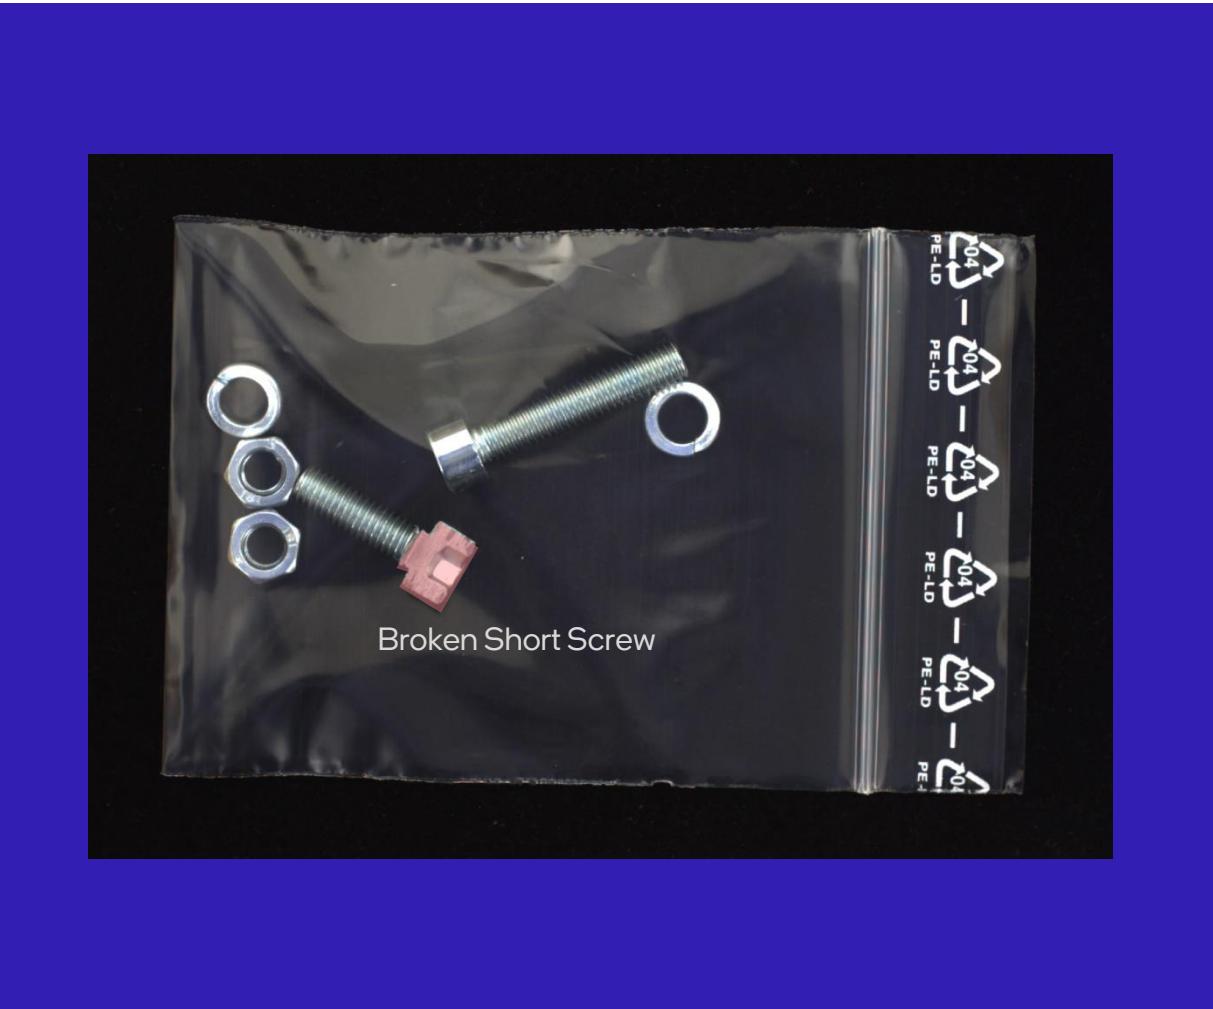
What / Where are the defects?



Where are the defects?



Where are the defects?



The Leaderboard

F1-Max^{1,2} will be the main evaluation metric for the challenge

- Category 1 final score is the harmonic mean of the Image F1 and Pixel F1 scores.
- Category 2 final score is the Image F1 score.

Category I – Adapt & Detect

Place	Entrant	Project	Image Score	Pixel Score	Final Score
1	Babar Hussain, Qiang Liu, Dahai Yu	ARNet for Robust Anomaly Detection	0.966	0.656	0.811
2	Canhui Tang, M. Callman, Sanping Zhou	CanhuiTang_submission_v2	0.947	0.598	0.773
3	Yukino Tsuzuki	Ensemble PatchCore	0.944	0.450	0.697

Category II – VLM Anomaly Challenge

Place	Entrant	Project	AUFC	Final Score
1	Zhaopeng Gu	AnomalyMoE	0.8184	0.8184
2	Nestor Bao	RJVoyagers_2	0.7958	0.7958
3	saidinesh pola	Category-2 VAND2.0 fewshot classification	0.7867	0.7867

¹<https://www.hackster.io/contests/openvino2024/>

²<https://sites.google.com/view/vand-2-0-cvpr-2024/challenge>

Winners



The Prizes



Grand Prize for Each Category

- The newest Intel AI PC: a laptop perfect for on-the-go AI boosting performance.
- An invitation to present their solutions during our workshop!



Second Place Prize for Each Category

- An Intel® ARC™ Discrete Graphics Card: Accelerate your system's performance.

Challenge Special Thanks



Dick Ameln, MSc

AI Research
Engineer/Scientist
Utrecht, NL



Ashwin Vaidya,
MSc

AI Research
Engineer/Scientist
Groningen, NL

Don't forget to install Anomalib and OpenVINO
for your next Anomaly detection project



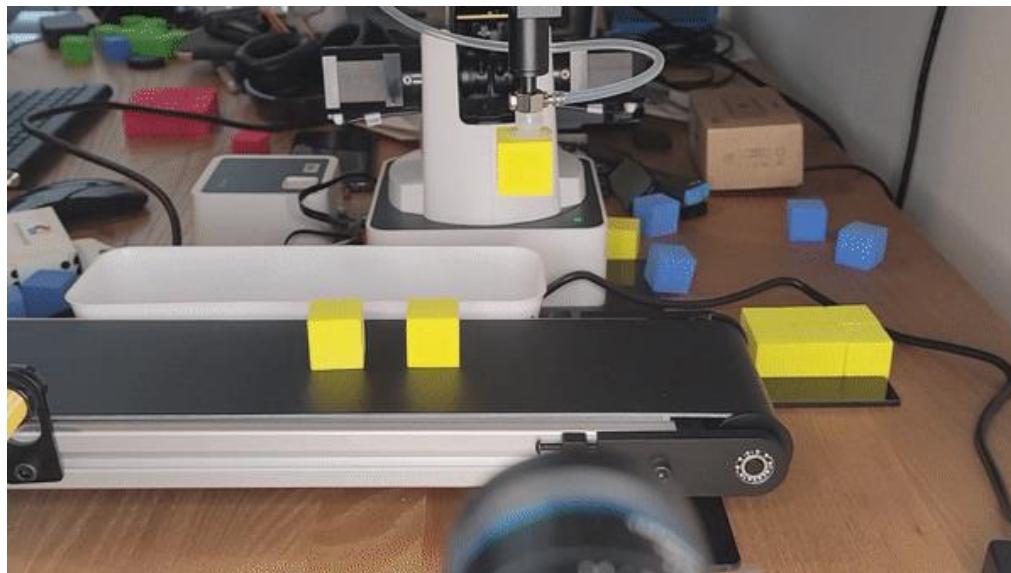
Anomalib Repo



OpenVINO Notebooks

Summary of Defect Detection with Anomalib

Solution Demo



Solution Flow

1. Training
2. Testing/Validating
3. Optimization
4. Deployment



Download from
our GitHub repo

Solution Highlights

- Unsupervised learning for imbalanced data set (PaDiM model)
- Optimized inference performance (OpenVINO)

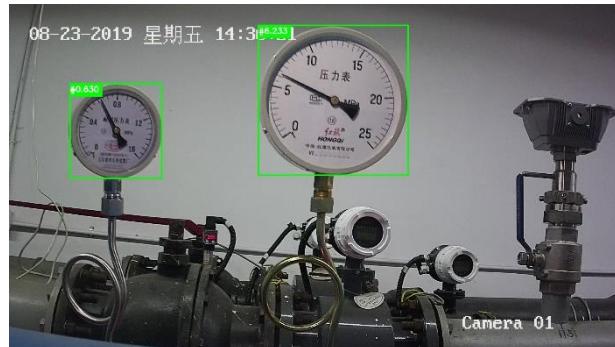
OpenVINO Features

- Reduce model size and improve inference speed while maintaining accuracy (quantization)
- Deploy to various intel and third-party hardware accelerators with minimal code changes

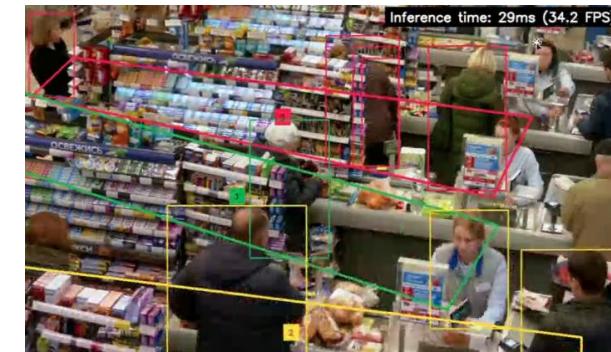
Anomalib Features

- Ready-to-use deep learning anomaly detection algorithms and benchmark datasets.
- PyTorch Lightning based model implementations to reduce boilerplate code and limit the implementation efforts to the bare essentials.
- A set of inference tools for quick and easy deployment of the standard or custom anomaly detection models.

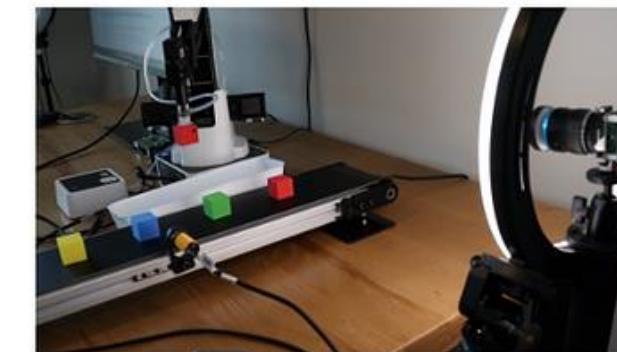
Build real-world AI solutions with Edge AI Reference Kits



Smart Meter
Scanning with
OpenVINO

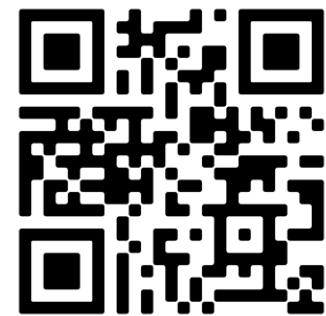


Intelligence Queue
Management with
OpenVINO

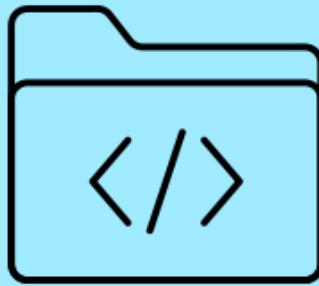


Defect Detection with
Anomalib and
OpenVINO

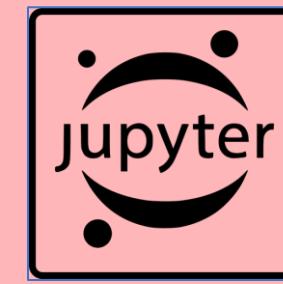
Explore more



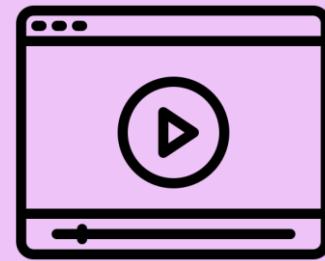
Jumpstart with rich materials and resources



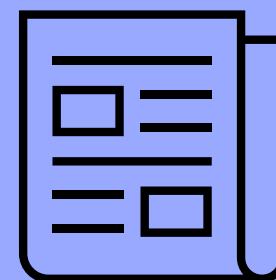
Source Codes



Jupyter Notebooks



Videos



Articles

Follow our AI software evangelists around the world



Ria Cheruvu
USA
(Phoenix, Arizona)
MSc in Data Science
Exp. 5+ years as
AI SW Lead Architect
and Research Engineer



Paula Ramos
USA
(Raleigh, North Carolina)
PhD in Engineering
Exp. 17+ years
as researcher



Dmitriy Pastushenkov
EMEA
(Karlsruhe, Germany)
MSc in
Computer Science
Exp. 17+ years as Software
Engineer and Architect



Adrian Boguszewski
EMEA
(Swindon, United Kingdom)
MS in Computer Science
Exp. 5+ years as
DL Engineer



Anisha Udayakumar
APJ
(Bengaluru, India)
BTech in Civil Engineering
Exp. 5+ years as
Innovation Consultant



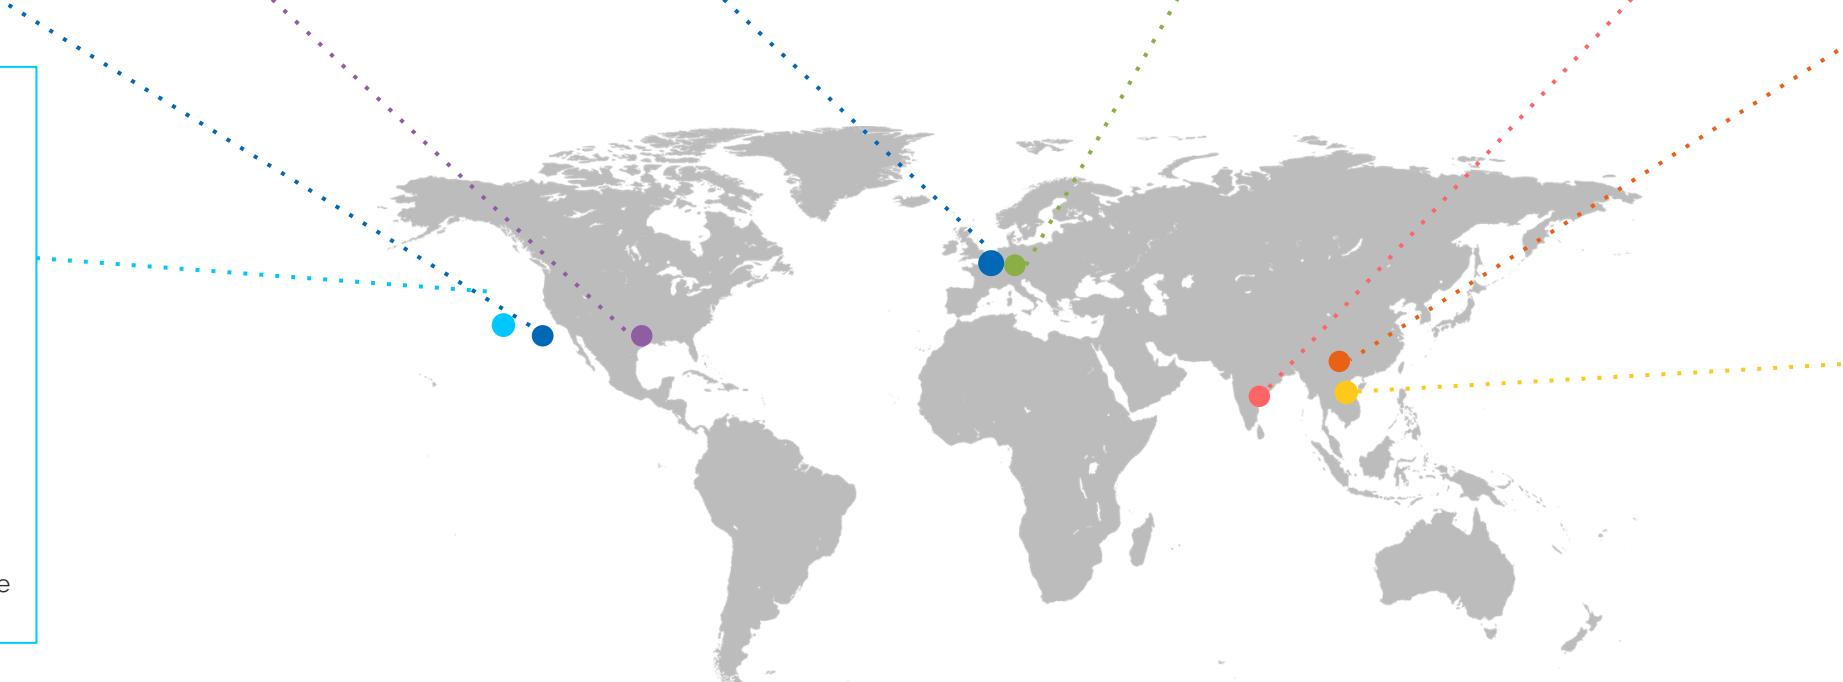
Zhuo Wu
PRC
(Shanghai, China)
PhD in Electronics
Exp. 15+ years as
researcher and professor



Raymond Lo
Global Lead
USA
(Santa Clara, California)
PhD Engineering
Exp. 9+ as Entrepreneur Executive
and Evangelist



Ethan Yang
PRC
(Shanghai, China)
MSc in Electronics
Exp 5+ years as
AI Solution Engineer



intel.
innovation

AI: The New Age Solving the World's Toughest Challenges, Together.

Calling All Developers & Technologists!

From front-end, web, app devs to back-end, full-stack, database & DevOps to data scientists & researchers, and more:
Learn, collaborate, and solve at Intel Innovation –
an event for developers by developers.

- ✓ Hear from leading industry luminaries, technologists & start-up entrepreneurs in the field of AI.
- ✓ Get the latest AI development tools, hands-on experience & join on-site Hackathons to optimize your AI code & workflows.
- ✓ Learn the breadth of future technology advancements in AI through keynotes, sessions, birds of a feathers, and hands-on labs.
- ✓ Share unique ideas and perspectives and collaborate with your peers.

Save the Date:
September 24-25, 2024
San Jose Convention Center, CA



Opt-In for Early Access
When Registration Opens!
www.intel.com/innovation

intel



Thank You



Paula Ramos

linkedin.com/in/paula-ramos-41097319/

Paula.Ramos@intel.com



Samet Akcay

linkedin.com/in/sametakcay

Samet.Akcay@intel.com



Notices and Disclaimers

Performance varies by use, configuration and other factors. Learn more at www.intel.com/PerformanceIndex.

Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See backup for configuration details.

No product or component can be absolutely secure.

Intel technologies may require enabled hardware, software or service activation.

Your costs and results may vary.

Intel is committed to respecting human rights and avoiding complicity in human rights abuses. See Intel's [Global Human Rights Principles](#). Intel's products and software are intended only to be used in applications that do not cause or contribute to a violation of an internationally recognized human right.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.